

Announcements

- ▶ another quiz/poll this morning
- ▶ posting sample solution to Lab01
- ▶ UNIX bootcamp this coming weekend
- ▶ CSSS election
 - more information at <http://csss.usask.ca>

Quote of the Day

- ▶ If the designers of X-Windows built cars, there would be no fewer than five steering wheels hidden about the cockpit, none of which followed the same principles -- but you'd be able to shift gears with your car stereo. Useful feature, that.
 - Marcus J. Ranum

UNIX File System Fundamentals

Noteworthy Directories

/	Root directory for the entire file system
./	Current directory
~/	Your home (login) directory
~user/	Home directory of given user
..	Parent directory

Noteworthy Directories

- ▶ `.` / used to signify executing the named file in the current working directory as a command

example: `./mycmd`

**typical way you will execute your command “./a.out”
it will just invoke which ever file**

File Conventions

- ▶ file “extensions” (c.f. Windows) are not required, but are useful
- ▶ for convenience don’t use spaces in filenames; instead try
 - `my_file_name`
 - `myFileName`

Basic File-Oriented Commands

- ▶ `cd`
- ▶ `pwd`
- ▶ `ls`
- ▶ `cp`
- ▶ `rm`
- ▶ `mv`
- ▶ `mkdir`, `rmdir`

File Permissions

- ▶ Files have three basic permissions
 - Read (r)
 - Write (w) (edit)
 - Execute/search (x)
- ▶ and three permission categories or levels of ownership
 - User (u)
 - Group (g)
 - Global/other (o)

File Permissions

- ▶ Controlled by a bit mask of the symbolic form
 - `rwXrwXrwX`
- ▶ each group of 3 bits corresponds to an ownership level
 - | | | | |
|------------------|------------------|------------------|--|
| user | group | other | stored as an octal number from 0-7
'rwx' = 111, 'r-' = 100 then change to octal |
| <code>rwX</code> | <code>rwX</code> | <code>rwX</code> | |
 - `r`, `w`, or `x` indicates permission on (a 1-bit)
 - `-` indicates permission off (a 0-bit)
- ▶ e.g. `ls -l temp`

File Permissions

► `chmod` “change permissions on a file”

- to change file permissions

- can use `chmod` symbolic-mode

- e.g. `chmod g+w temp` “turn on write for group”

- `chmod o-rwx temp`

- `chmod a+r temp`

- `chmod a+rwx temp`

- can use `chmod mask` where `mask` is a bit pattern in octal

- e.g. `chmod 764 temp`

u - owner/user

g - group

o - other

a - all

+ for add

- for taking away

Controlling Ownership

- ▶ `chgrp`
 - change group ownership
 - `usage chgrp groupname file`
- ▶ `chown` **you need to own the file in the first place**
 - `chown` owner of file
 - typically a restricted command for security and sys admin reasons
- ▶ note: typically on lab machines, students are restricted in ability to change file ownership

File Types

- ▶ basic types: ordinary files, directories
- ▶ other types
- ▶ how to determine the type of an ordinary file?
 - especially since extensions are not required
 - `file` command
 - makes use of “magic pattern” information, typically at the beginning of the file
 - `man magic`
 - example

Pattern-Matching in Filenames

- ▶ *****
 - match any number of any characters
 - ▶ **?**
 - match any one character
 - ▶ **[]**
 - match one of a set of characters
 - ▶ **others**
 - ▶ **N.B.: supported by shell, and not by the file system nor application program**
- Is temp*** (will match anything that starts with temp)
Is temp?.pdf (it will match temp and then any char.pdf)
Is *pdf (lists all pdf files)
Is temp[15].pdf (matches 1 char that comes from set [15], so any temp1...temp5.pdf)
Is .* (lists all files that start with '.')
Is -d .* (will stop at the directory, will not list contents. Will list all directories that start with '.')
- # all directories have '.' (list to itself) and '..' (list to its parent)**
- Is t???*.txt** (all names that start with a 't' then 3 char.txt)

More Commands Related to Files

- ▶ `more`, `less`
- ▶ `diff`, `cmp` comparing content of files (only on ASCII text), are these 2 files the same?
- ▶ `WC` counts things. wants to work on txt files
- ▶ `sort` sorts txt files
- ▶ `uniq` whether or not all the lines in the txt file are unique. Assumes its sorted
- ▶ `head`, `tail` gives begining or end of file
- ▶ `du` how much disk usage. -s means sum up all, -m means megabytes
- ▶ `df` how much free space. What files systems are mounted.
- ▶ `example`

More Commands Related to Files

- ▶ changing files between operating systems

- DOS to UNIX

- `unix2dos` and `dos2unix` on tuxworld

- `tr '\r' '\n'`

- `mttools` on tuxworld

Special Files

- ▶ `/dev/null` anything you write here, the OS will get rid of it
 - infinite sink
 - infinite source of end-of-file
 - uses
 - discarding output
 - terminating input

Working With Files

- ▶ Recall: “Everything in UNIX is a file”
 - useful hyperbole
- ▶ Every open file assigned a number called a *(file) descriptor*
 - small integer, starting at 0
 - think of it as a “pointer to an open file”
 - unique set per process

Standard Files

- ▶ 3 files automatically associated with every process, with every command invoked by the shell
- ▶ *stdin* - the input stream
- ▶ *stdout* - the output stream
- ▶ *stderr* - the error (output) stream
- ▶ in C, we will see these as `stdin`, `stdout`, and `stderr` defined by the `stdio` library
- ▶ in C++, we have objects `cin`, `cout`, `cerr` described as part of `cstdio`

stdin

- ▶ file descriptor 0
- ▶ default binding to the keyboard
- ▶ by default, inputs to a program are read from *stdin*

stdout

- ▶ file descriptor 1
- ▶ default binding to the terminal or display
- ▶ by default, program output is written to *stdout*

stderr

- ▶ file descriptor 2
- ▶ default binding to the terminal or display
- ▶ by default, any error or warning messages are (supposed to be) written to *stderr*
 - some programs do not comply; watch out!
 - in our code, error and warning messages should be sent to *stderr*, not *stdout*

Redirection Revisited

- ▶ earlier, saw examples like
`cat < source > destination`
- ▶ in `bash` can explicitly redirect file descriptors by preceding the ‘>’ or ‘<’ with the file descriptor number
 - e.g. above equivalent to
`cat 0< source 1> destination`
- ▶ saw `2> file` for redirection of *stderr* earlier

Redirection Revisited

- ▶ can also duplicate a file descriptor
 - two file descriptors will refer to the same file
 - full semantics beyond the scope of this course
- ▶ in `bash`, accomplished by adding ‘&’ to redirection operators
 - e.g. `>&n` instead of `>`

Redirection Revisited

- ▶ common use: to redirect both stdout and stderr to a single file

- e.g. in bash

```
prog >log 2>&1
```

- ▶ evaluation is left-to-right
- ▶ note: order of evaluation is important!

```
prog 2>&1 >log
```

does something different.

What is a Process?


- ▶ one definition of a *process*
 - a thread of control in an address space
- ▶ recall:
 - a program may invoke several processes
 - a single process can run multiple programs

Basic Process Abstraction in UNIX

- ▶ processes exist in a hierarchy
- ▶ parent/child/sibling model
 - each process has a unique parent
 - processes can have multiple children
 - each child will be a sibling of the other children
- ▶ each process identified by a unique identifier, its *PID*

Basic Process Abstraction in UNIX

► example

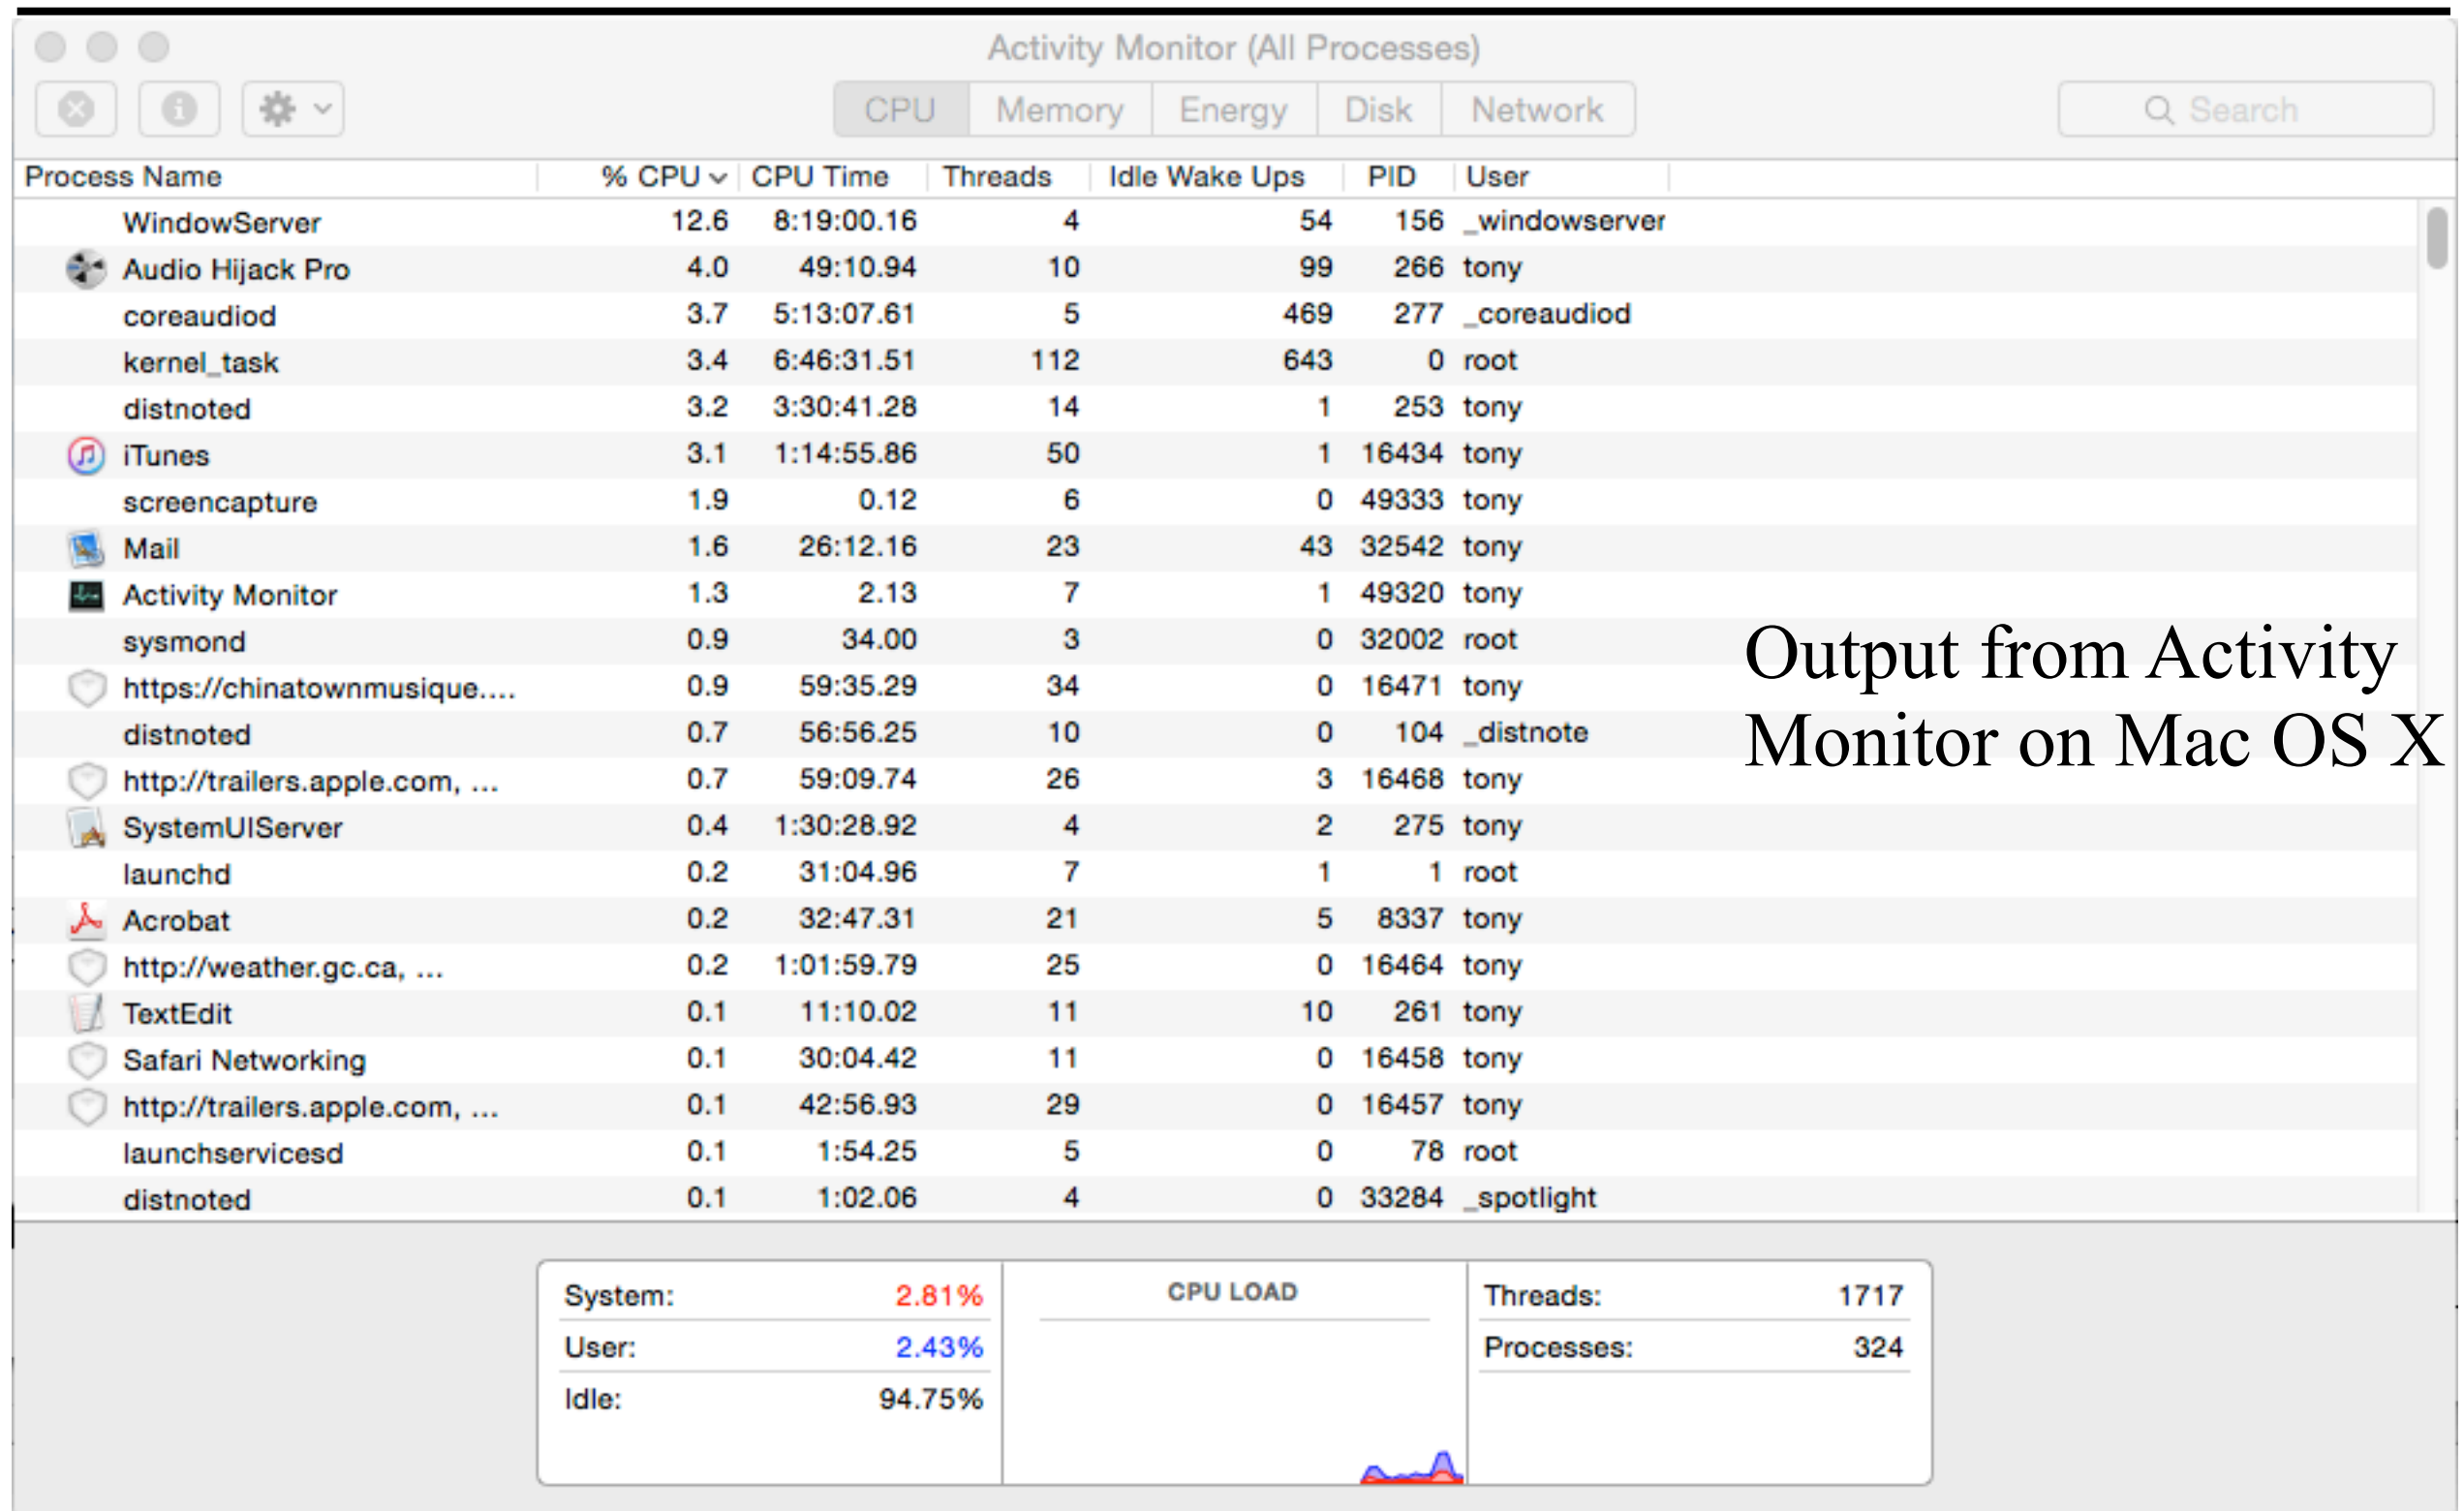


Basic Process Abstraction in UNIX

- abstraction in other operating systems is similar

COMMAND	PID	USER	TIME	%KER	%USE	PRI	RSS	SWAP	%MEM	THRD	%CPU
top	2620	administra	0:00	100	0	8	2076	676	0.20	1	40.00
lsass	672	SYSTEM	1h42	28	71	9	80240	77764	7.66	56	0.20
mstsc	2128	administra	25:12	48	51	8	5928	8504	0.57	10	0.12
cmd	1528	administra	0:00	71	28	8	1512	1424	0.14	1	0.05
services	660	SYSTEM	6:56	46	53	9	136580	4372	13.03	20	0.01
dns	1976	SYSTEM	5:24	53	46	8	7428	9064	0.71	14	0.01
mmc	2712	administra	0:08	62	37	8	16464	9108	1.57	5	0.01
svchost	1340	SYSTEM	4:47	41	58	8	24116	17340	2.30	41	0.01
winlogon	2884	SYSTEM	0:05	16	83	13	6412	6028	0.61	15	0.01
winlogon	600	SYSTEM	3:27	57	42	13	4796	7116	0.46	22	0.01
perl	1644	administra	2:16	24	75	8	15720	9752	1.50	4	0.00
dfssvc	1944	SYSTEM	1:52	49	50	8	4724	1892	0.45	11	0.00
svchost	1180	-	1:32	64	35	8	3652	1340	0.35	10	0.00
explorer	3540	administra	1:26	79	20	8	18172	8588	1.73	8	0.00
spoolsv	1720	SYSTEM	1:10	34	65	8	7796	5196	0.74	17	0.00
csrss	1520	SYSTEM	0:01	65	34	13	3024	1076	0.29	11	0.00
explorer	424	administra	0:00	69	30	8	10800	6368	1.03	10	0.00

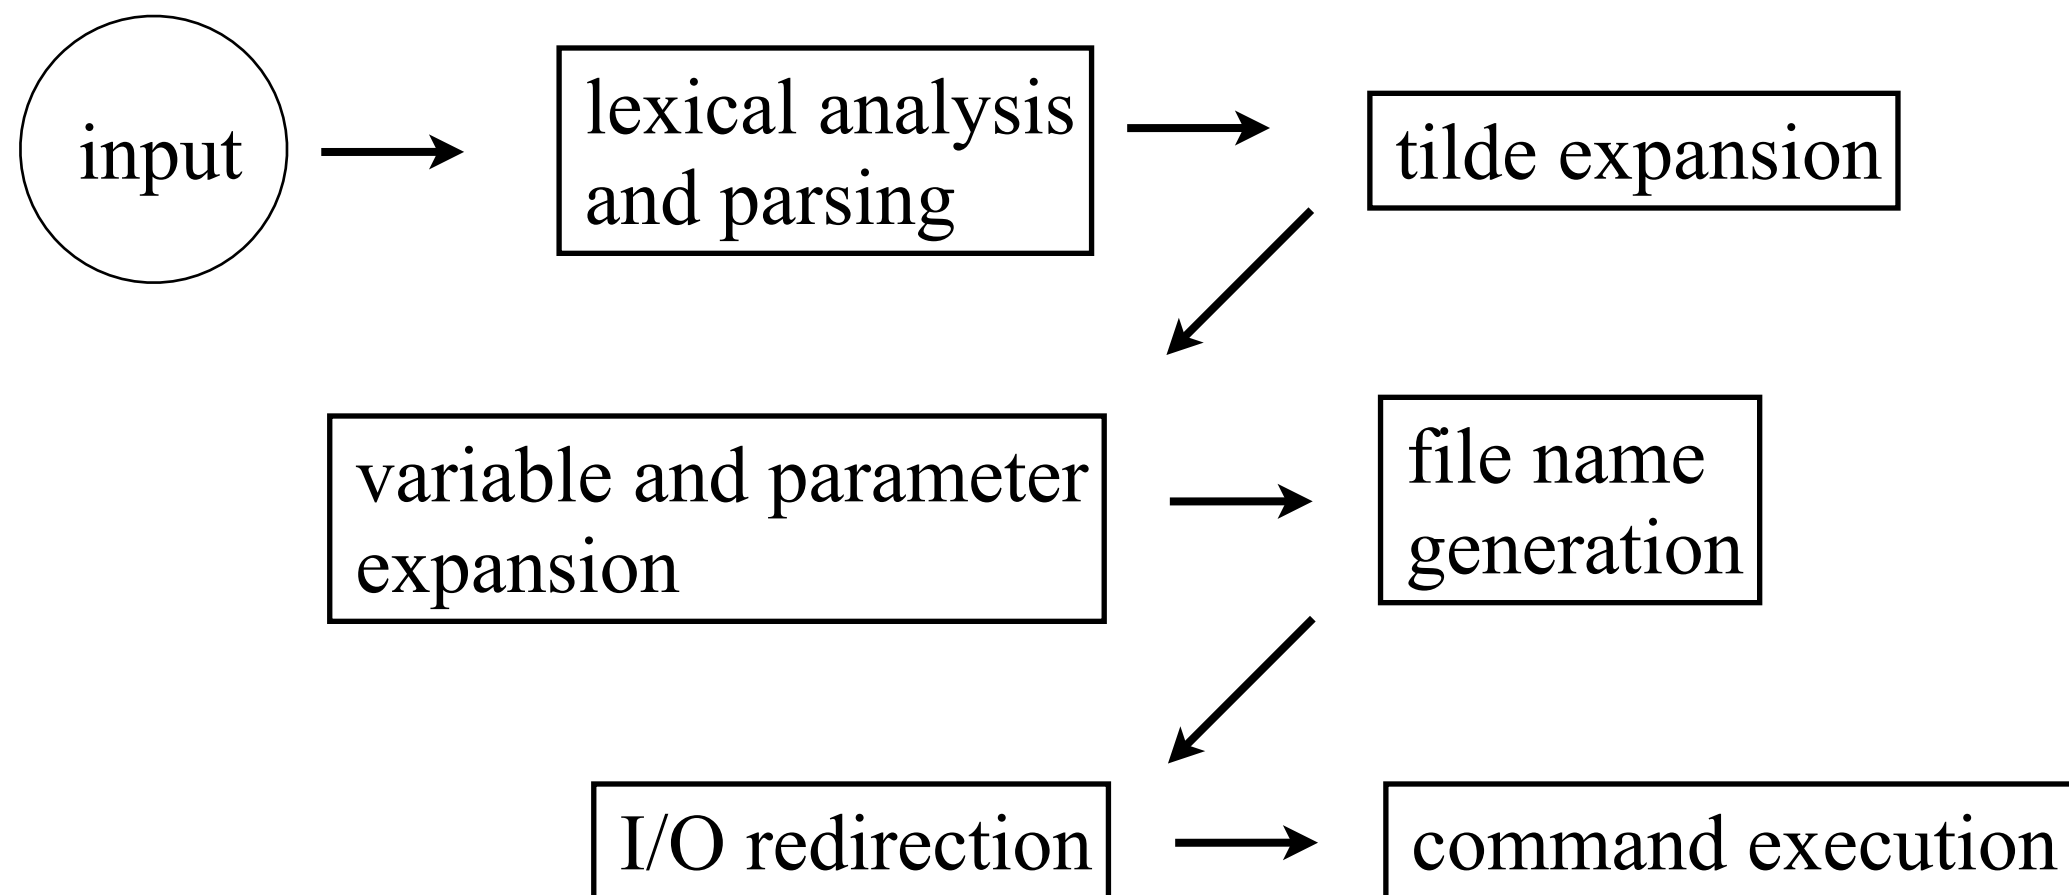
Basic Process Abstraction in UNIX



Output from Activity
Monitor on Mac OS X

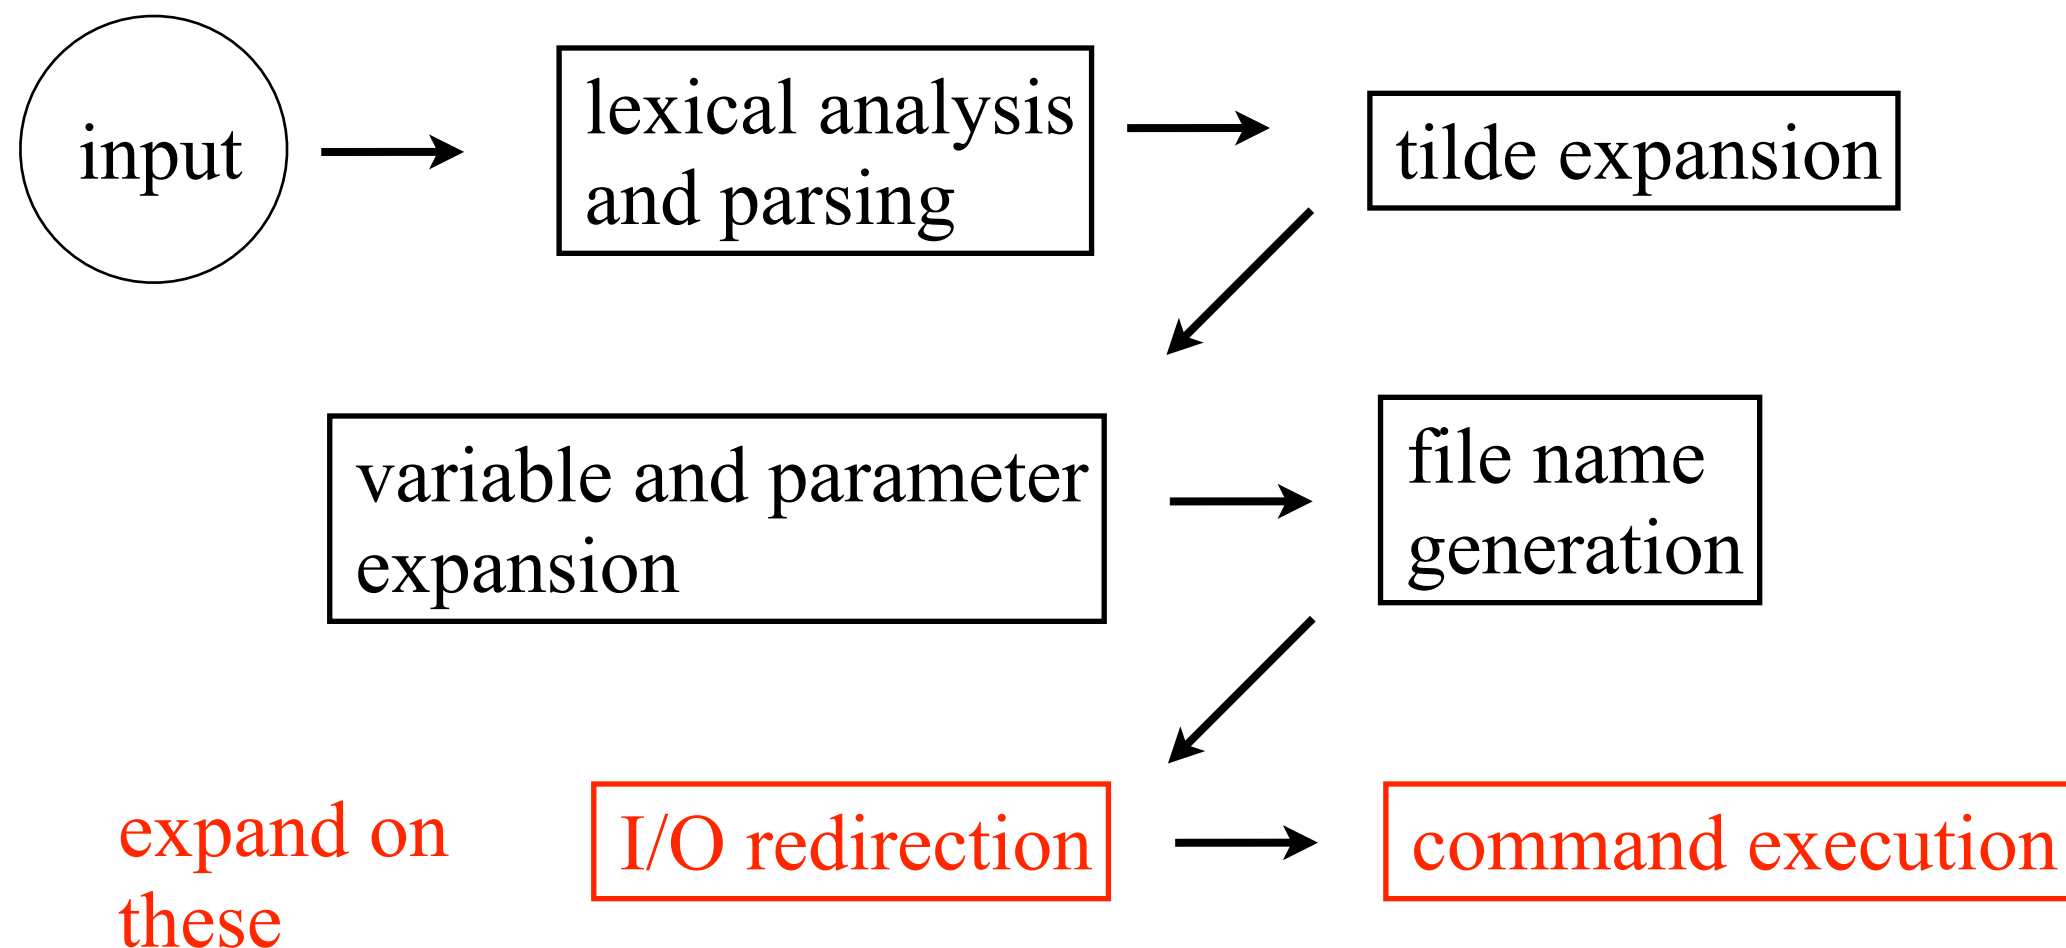
Basic Process Abstraction in UNIX

- ▶ process abstraction involved in executing a command from the shell
 - for simplicity many stages not shown



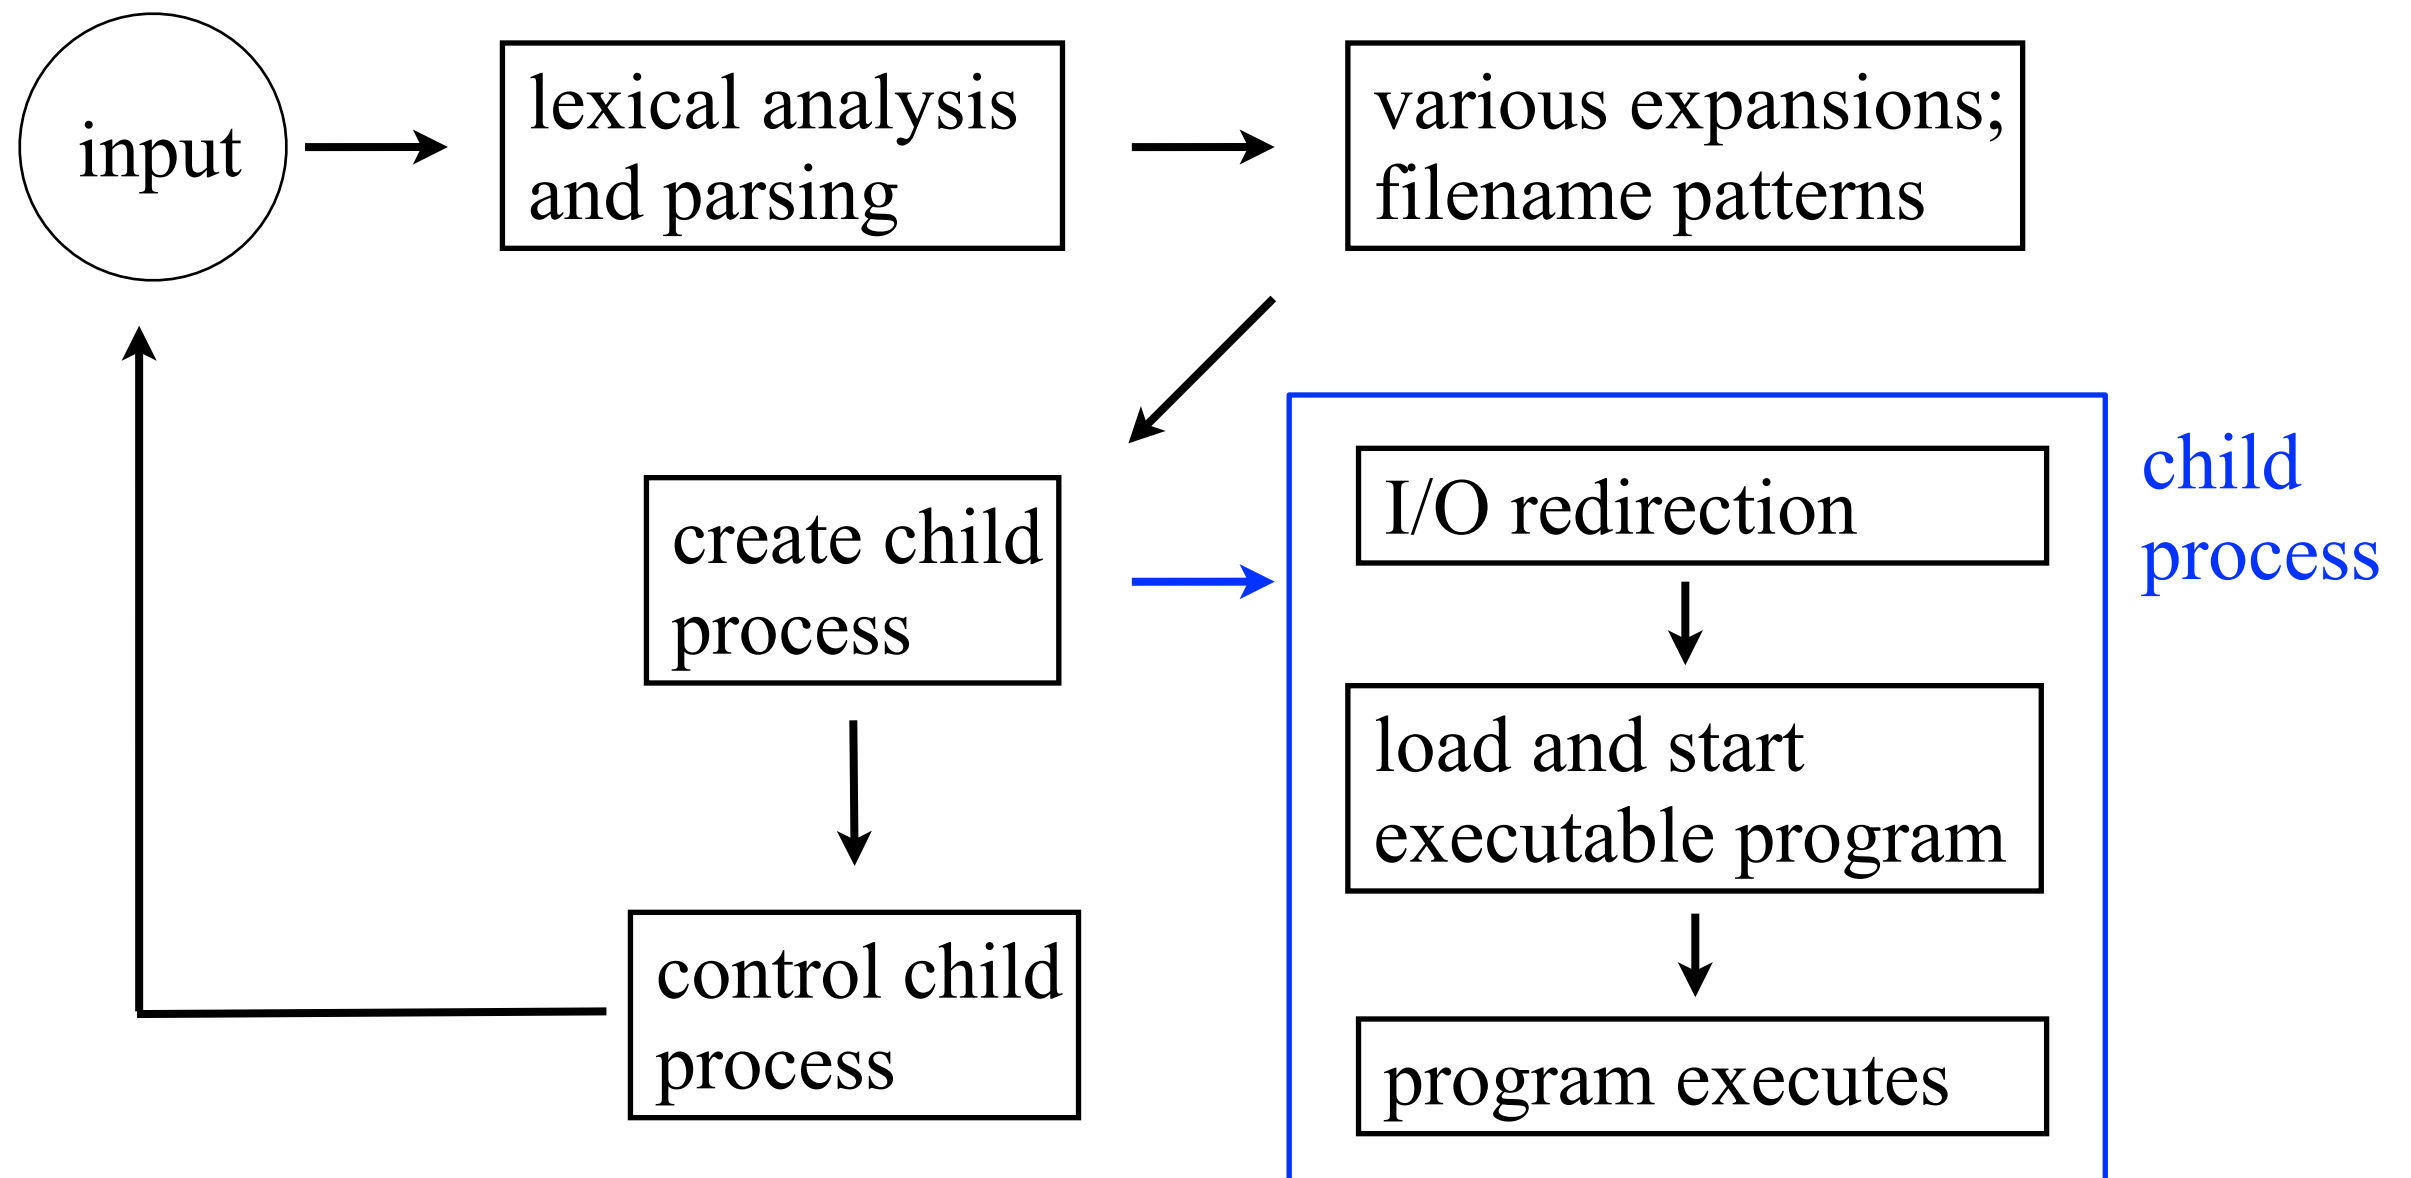
Basic Process Abstraction in UNIX

- ▶ process abstraction involved in executing a command from the shell
 - for simplicity many stages not shown



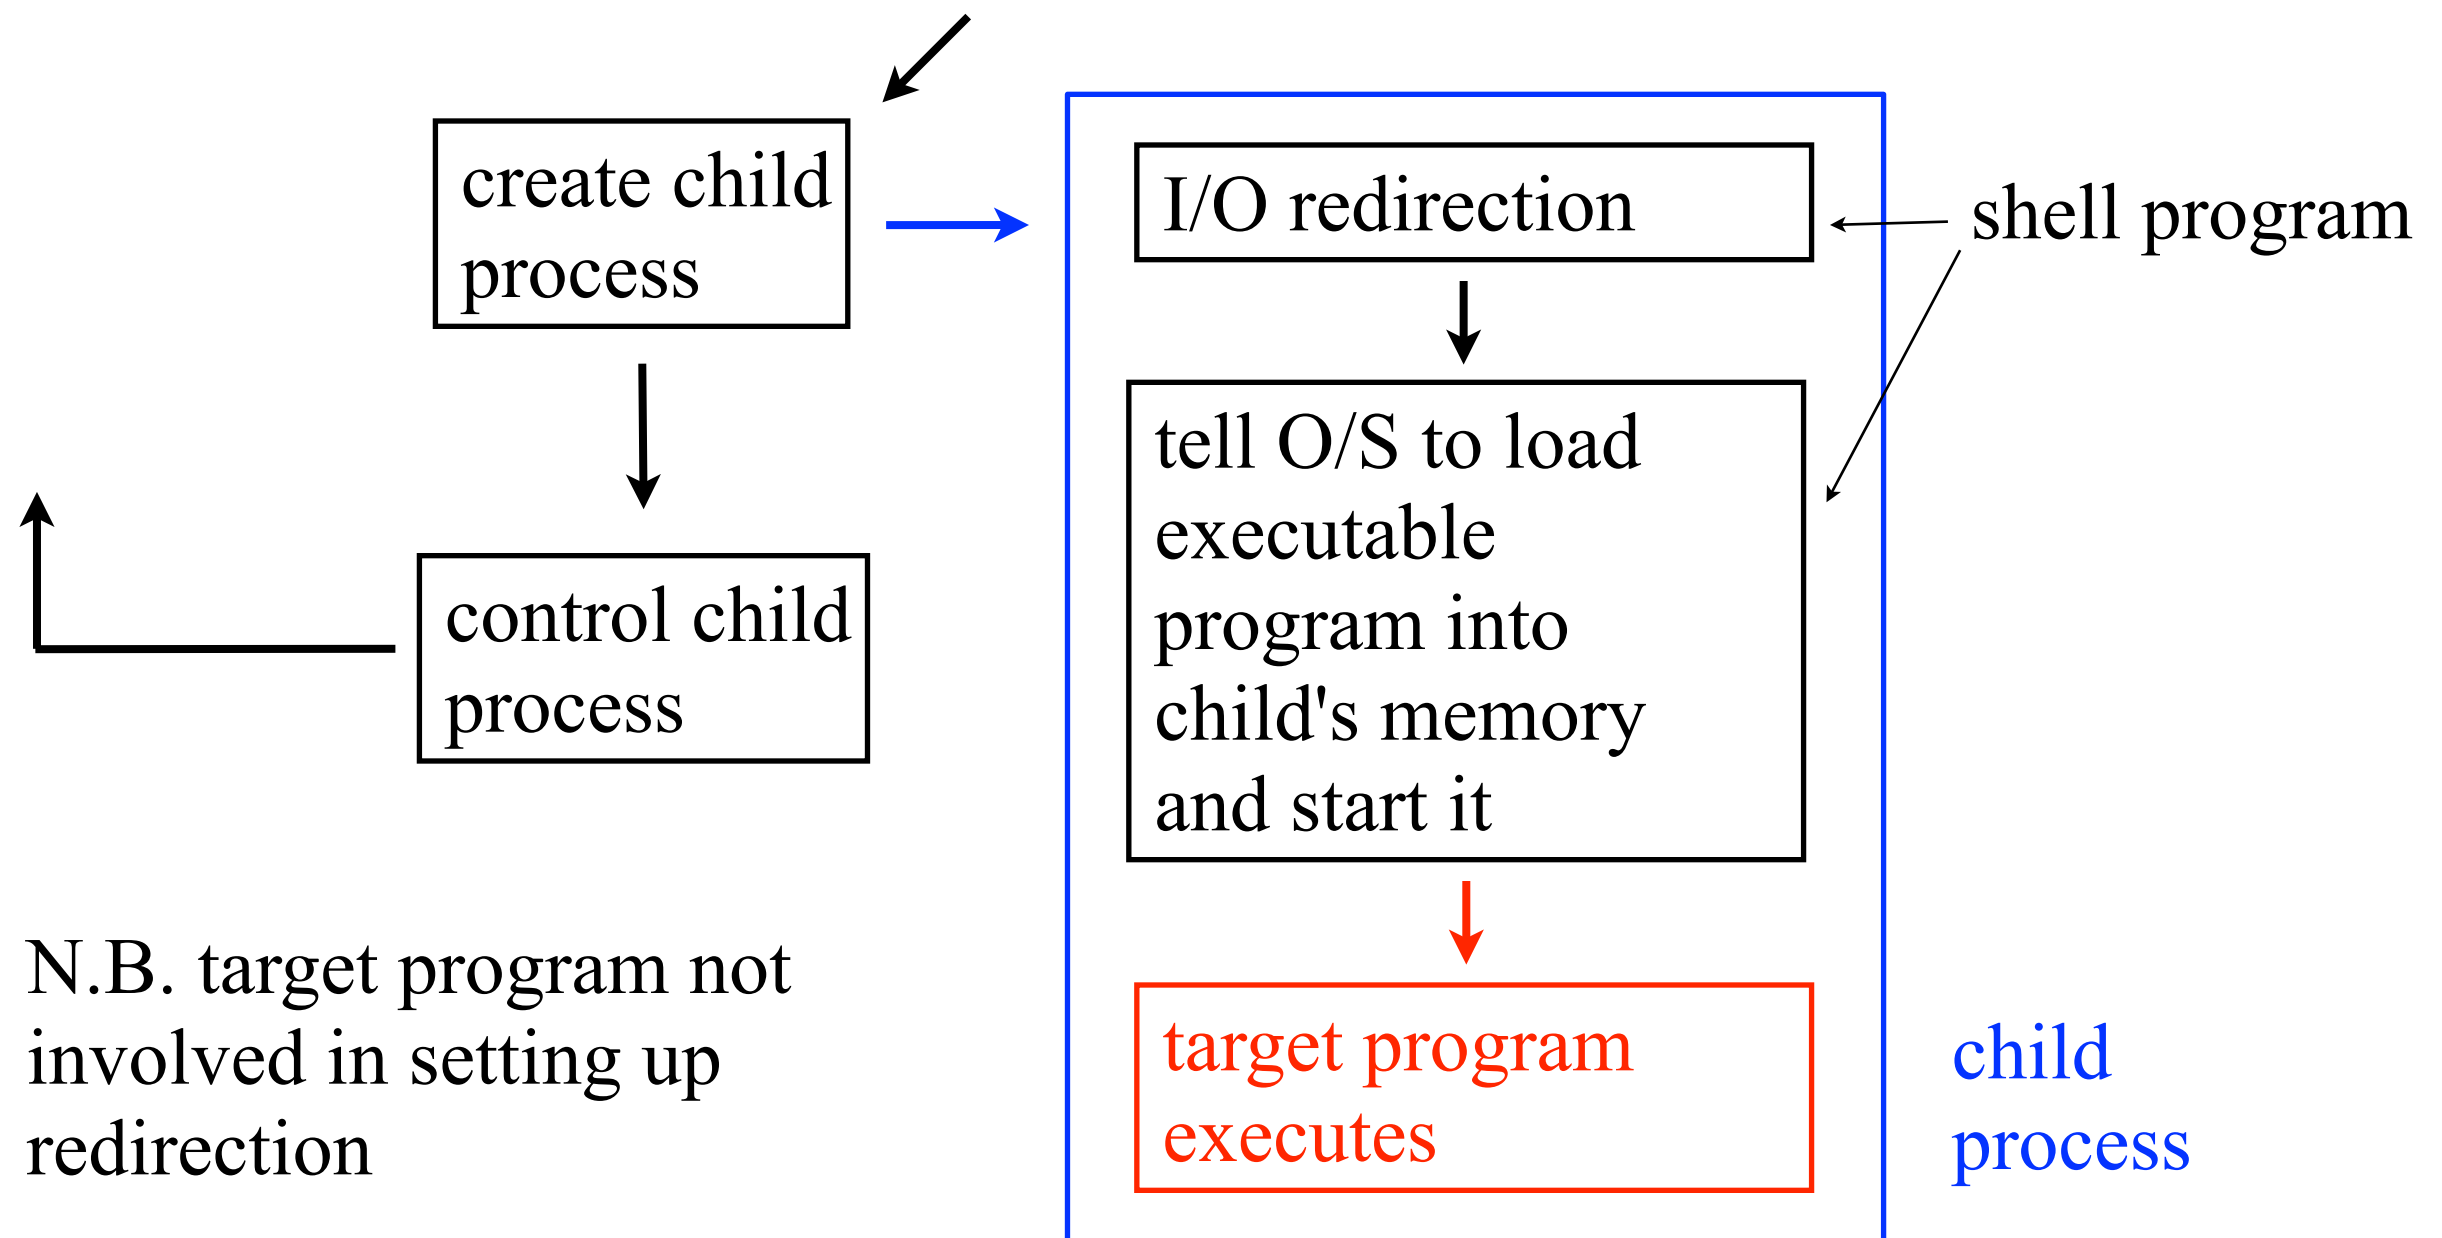
Basic Process Abstraction in UNIX

- ▶ process abstraction involved in executing a command from the shell



Basic Process Abstraction in UNIX

- ▶ process abstraction involved in executing a command from the shell



Commands Related to UNIX Processes

- ▶ list processes

- `ps`
- `pstree -h on tuxworld`
- `top`

- ▶ `uptime`

- ▶ `w` **and** `who`

- ▶ `exit` (built-in) and `^D` (end-of-file)

Commands Related to UNIX Processes

► eliminate processes

- `kill`
- `man 7 signal`
- signals generated by keyboard action: `SIGINT`, `SIGQUIT`
- useful signals for users: `SIGKILL`, `SIGTERM`
- `kill` built-in for `csch`, `/usr/bin/kill` **or** `/bin/kill` for `bash`
- `man 1 kill` **or** `info kill`

Processes and Jobs

- ▶ warning: UNIX shell specific definitions
- ▶ *foreground* process:
 - a process that is associated with user input
 - usually means “has control of the keyboard”
 - shell waits for its completion
- ▶ *background* process:
 - a process that executes whenever permitted by the OS
 - usually means “does not require user interaction”
 - shell does not wait for its completion

Processes and Jobs

▶ suspended process:

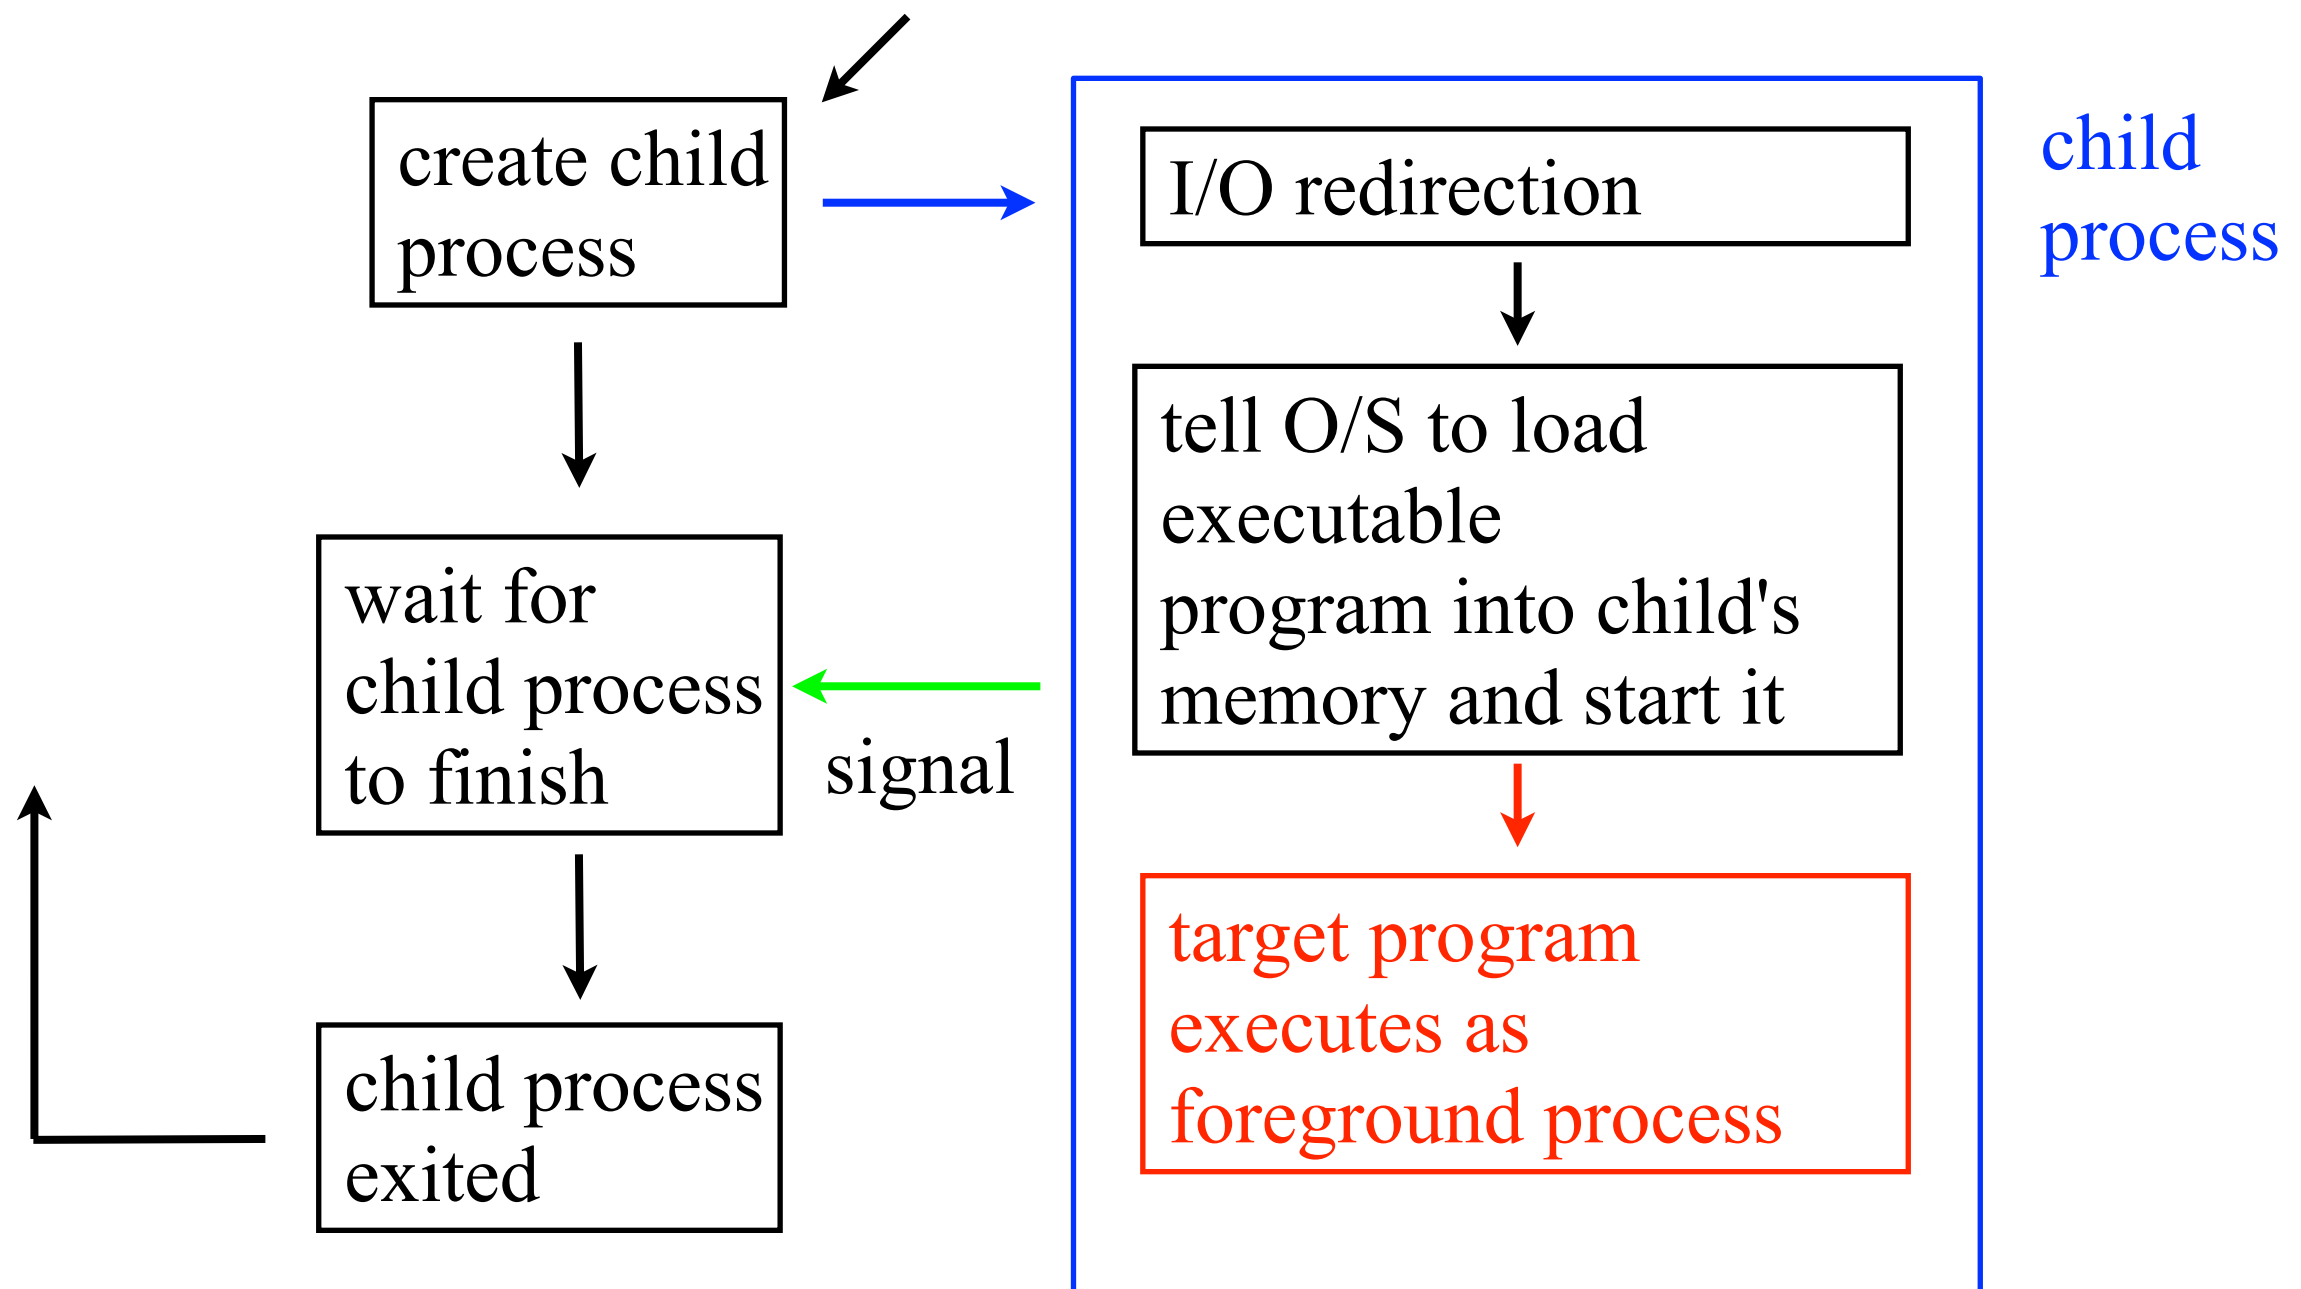
- a process that was executing, as permitted by the OS, but is now inactive
- usually means “was consuming computing resources, but is no longer doing so”. However, the process is still likely using memory resources (e.g. RAM)
- `ps -l` (LINUX) or `ps -av` (BSD)

▶ *job*:

- a suspended or background running process
- `jobs`

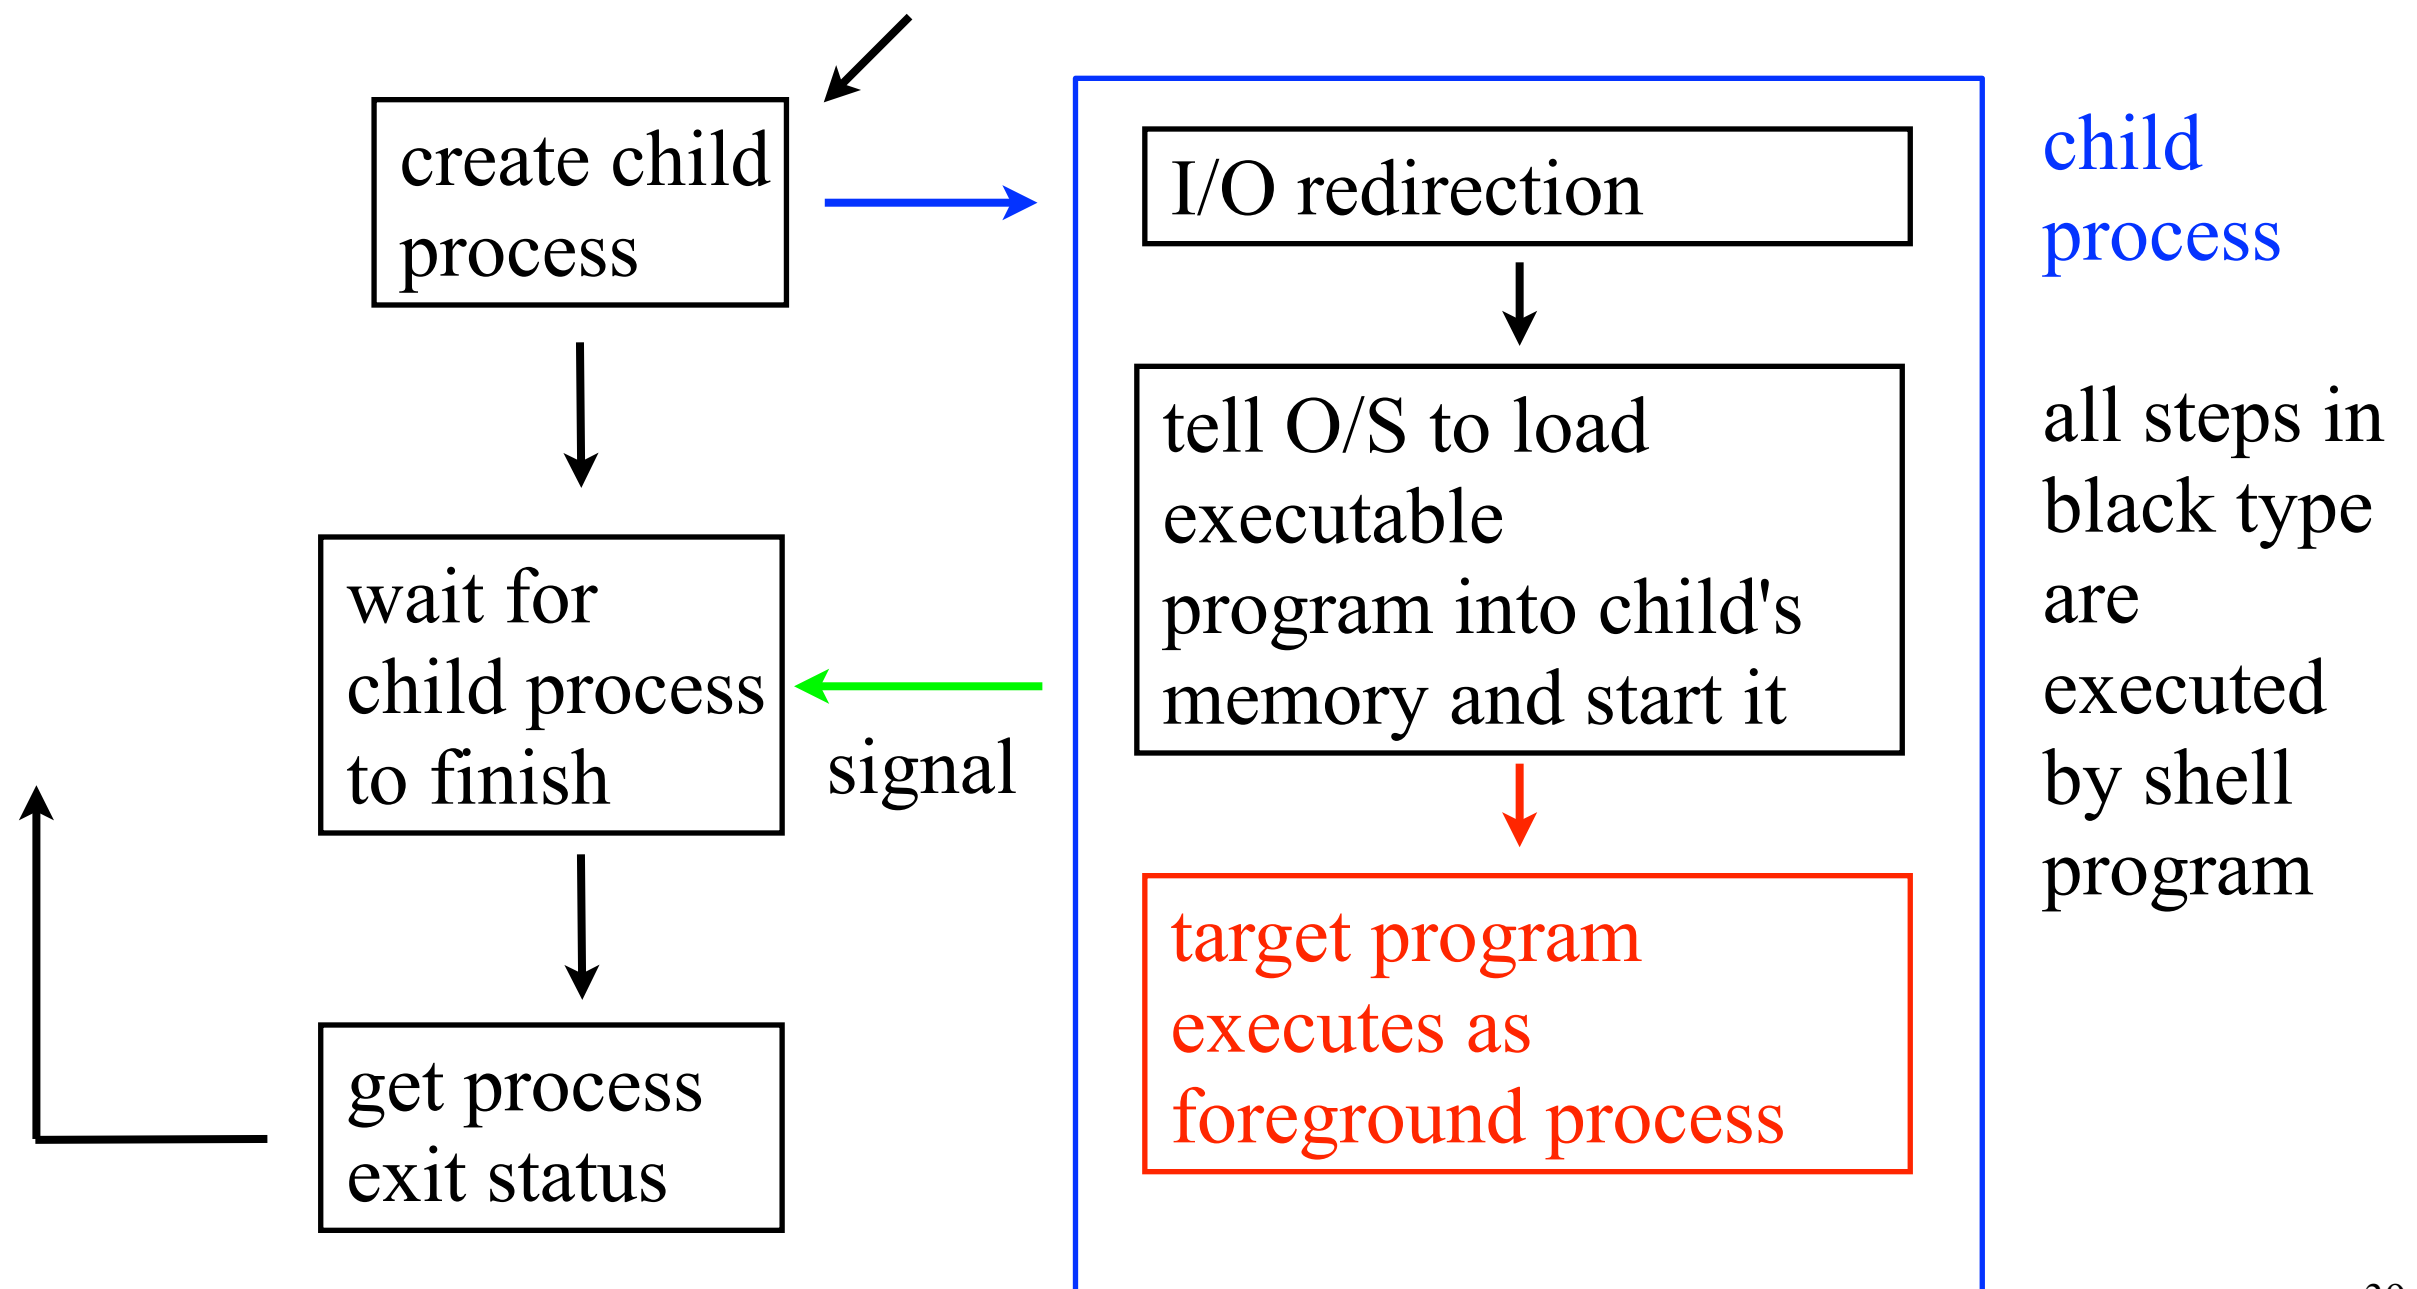
Basic Process Abstraction in UNIX

- ▶ process abstraction involved in executing a foreground command from the shell



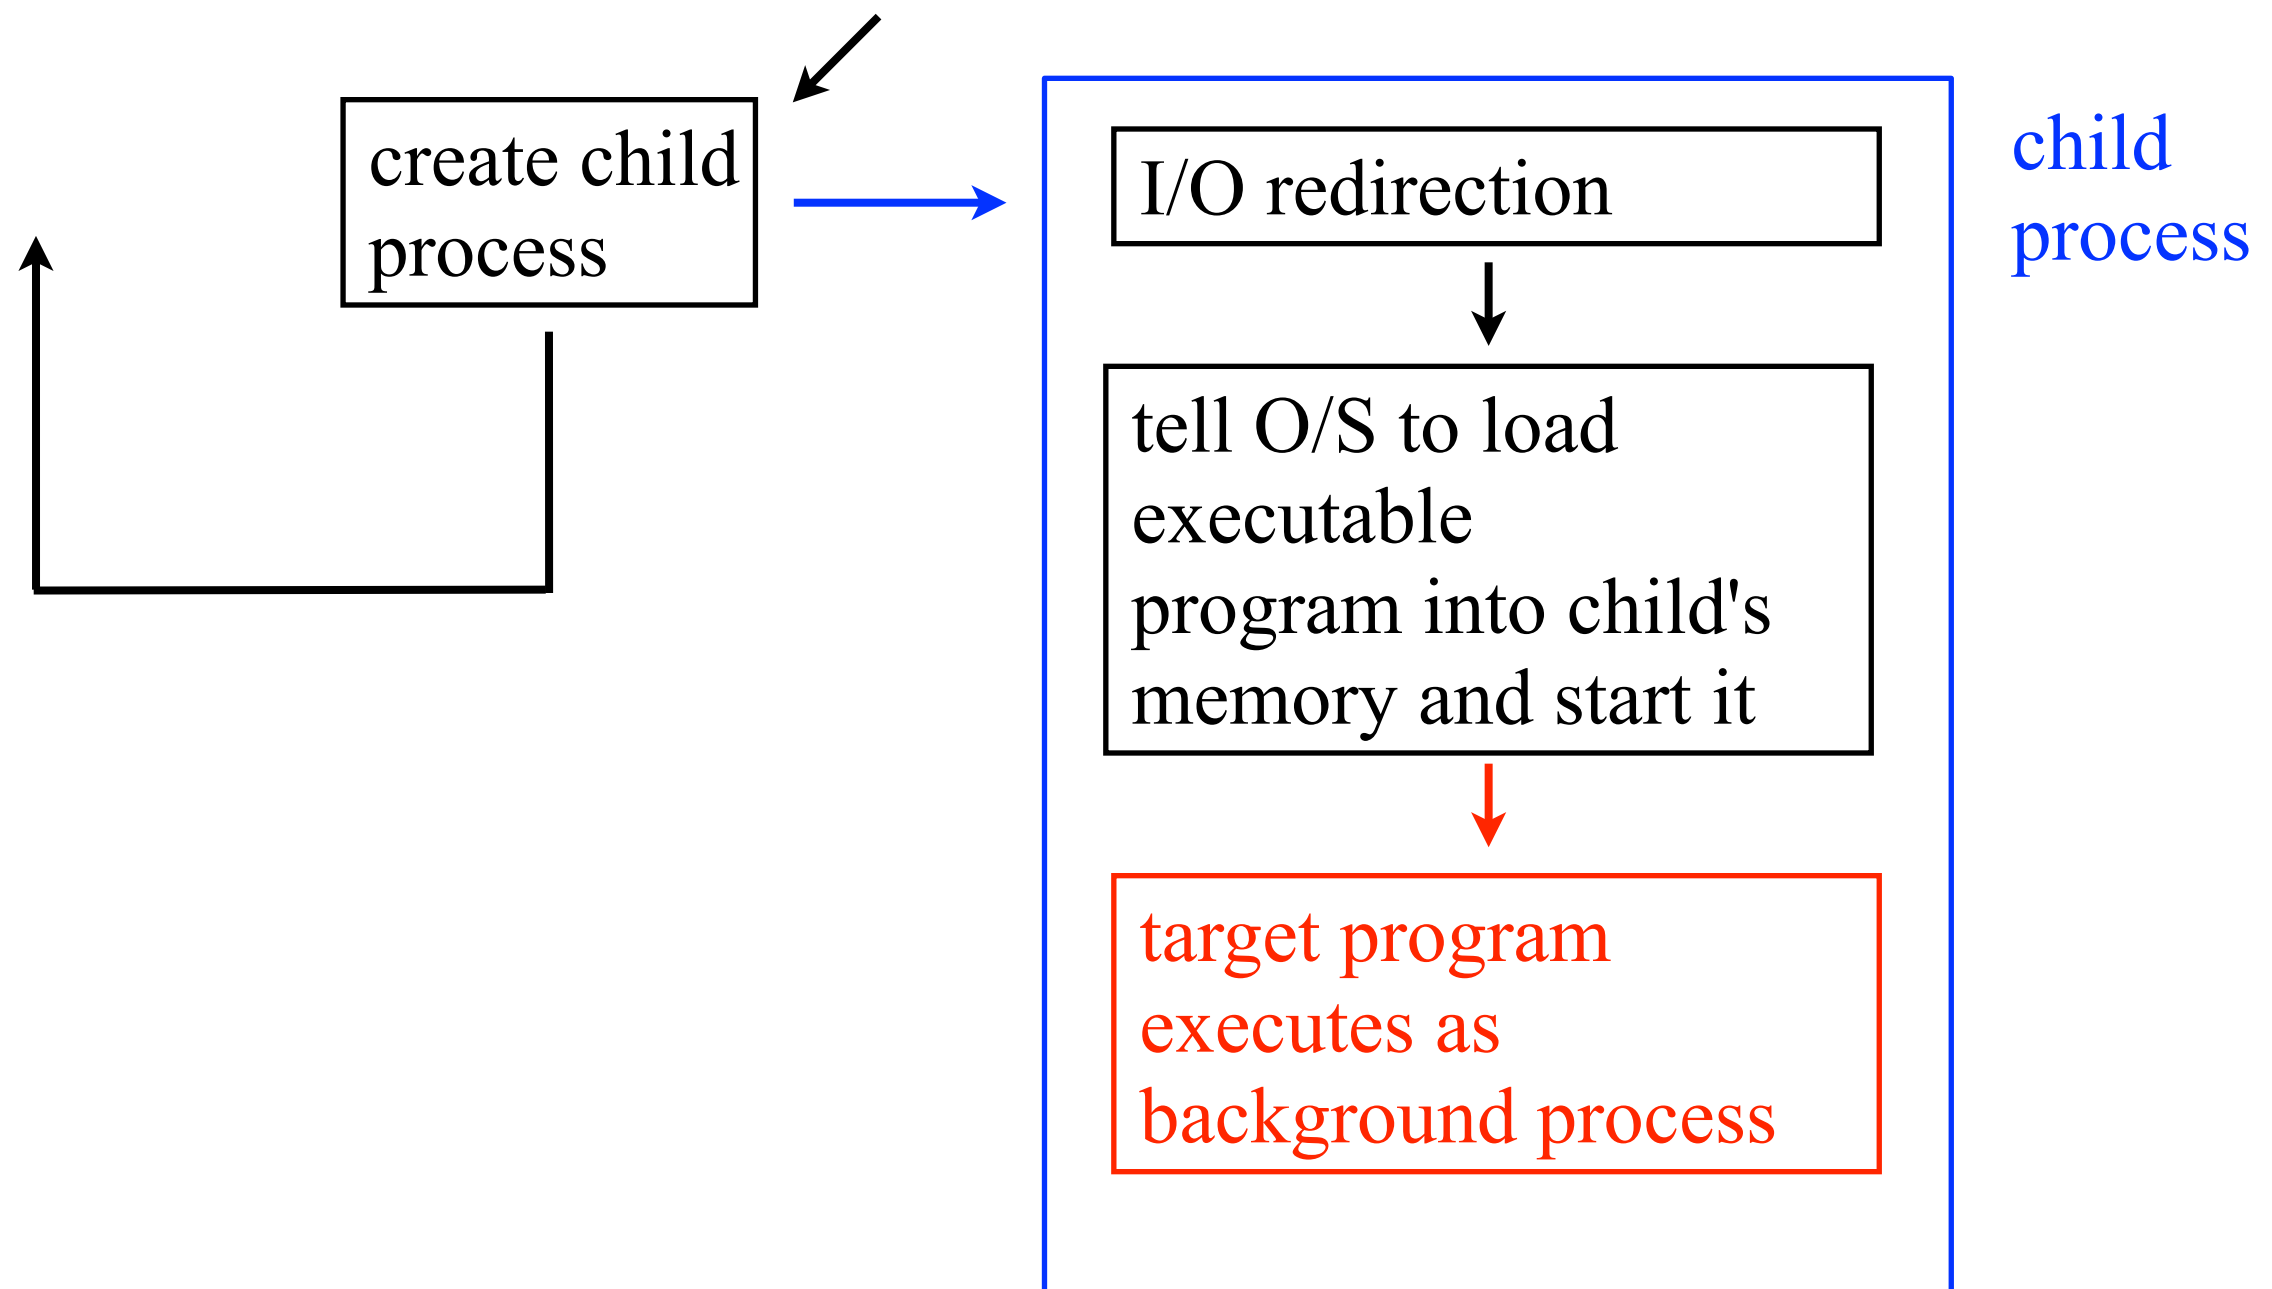
Basic Process Abstraction in UNIX

- ▶ process abstraction involved in executing a foreground command from the shell



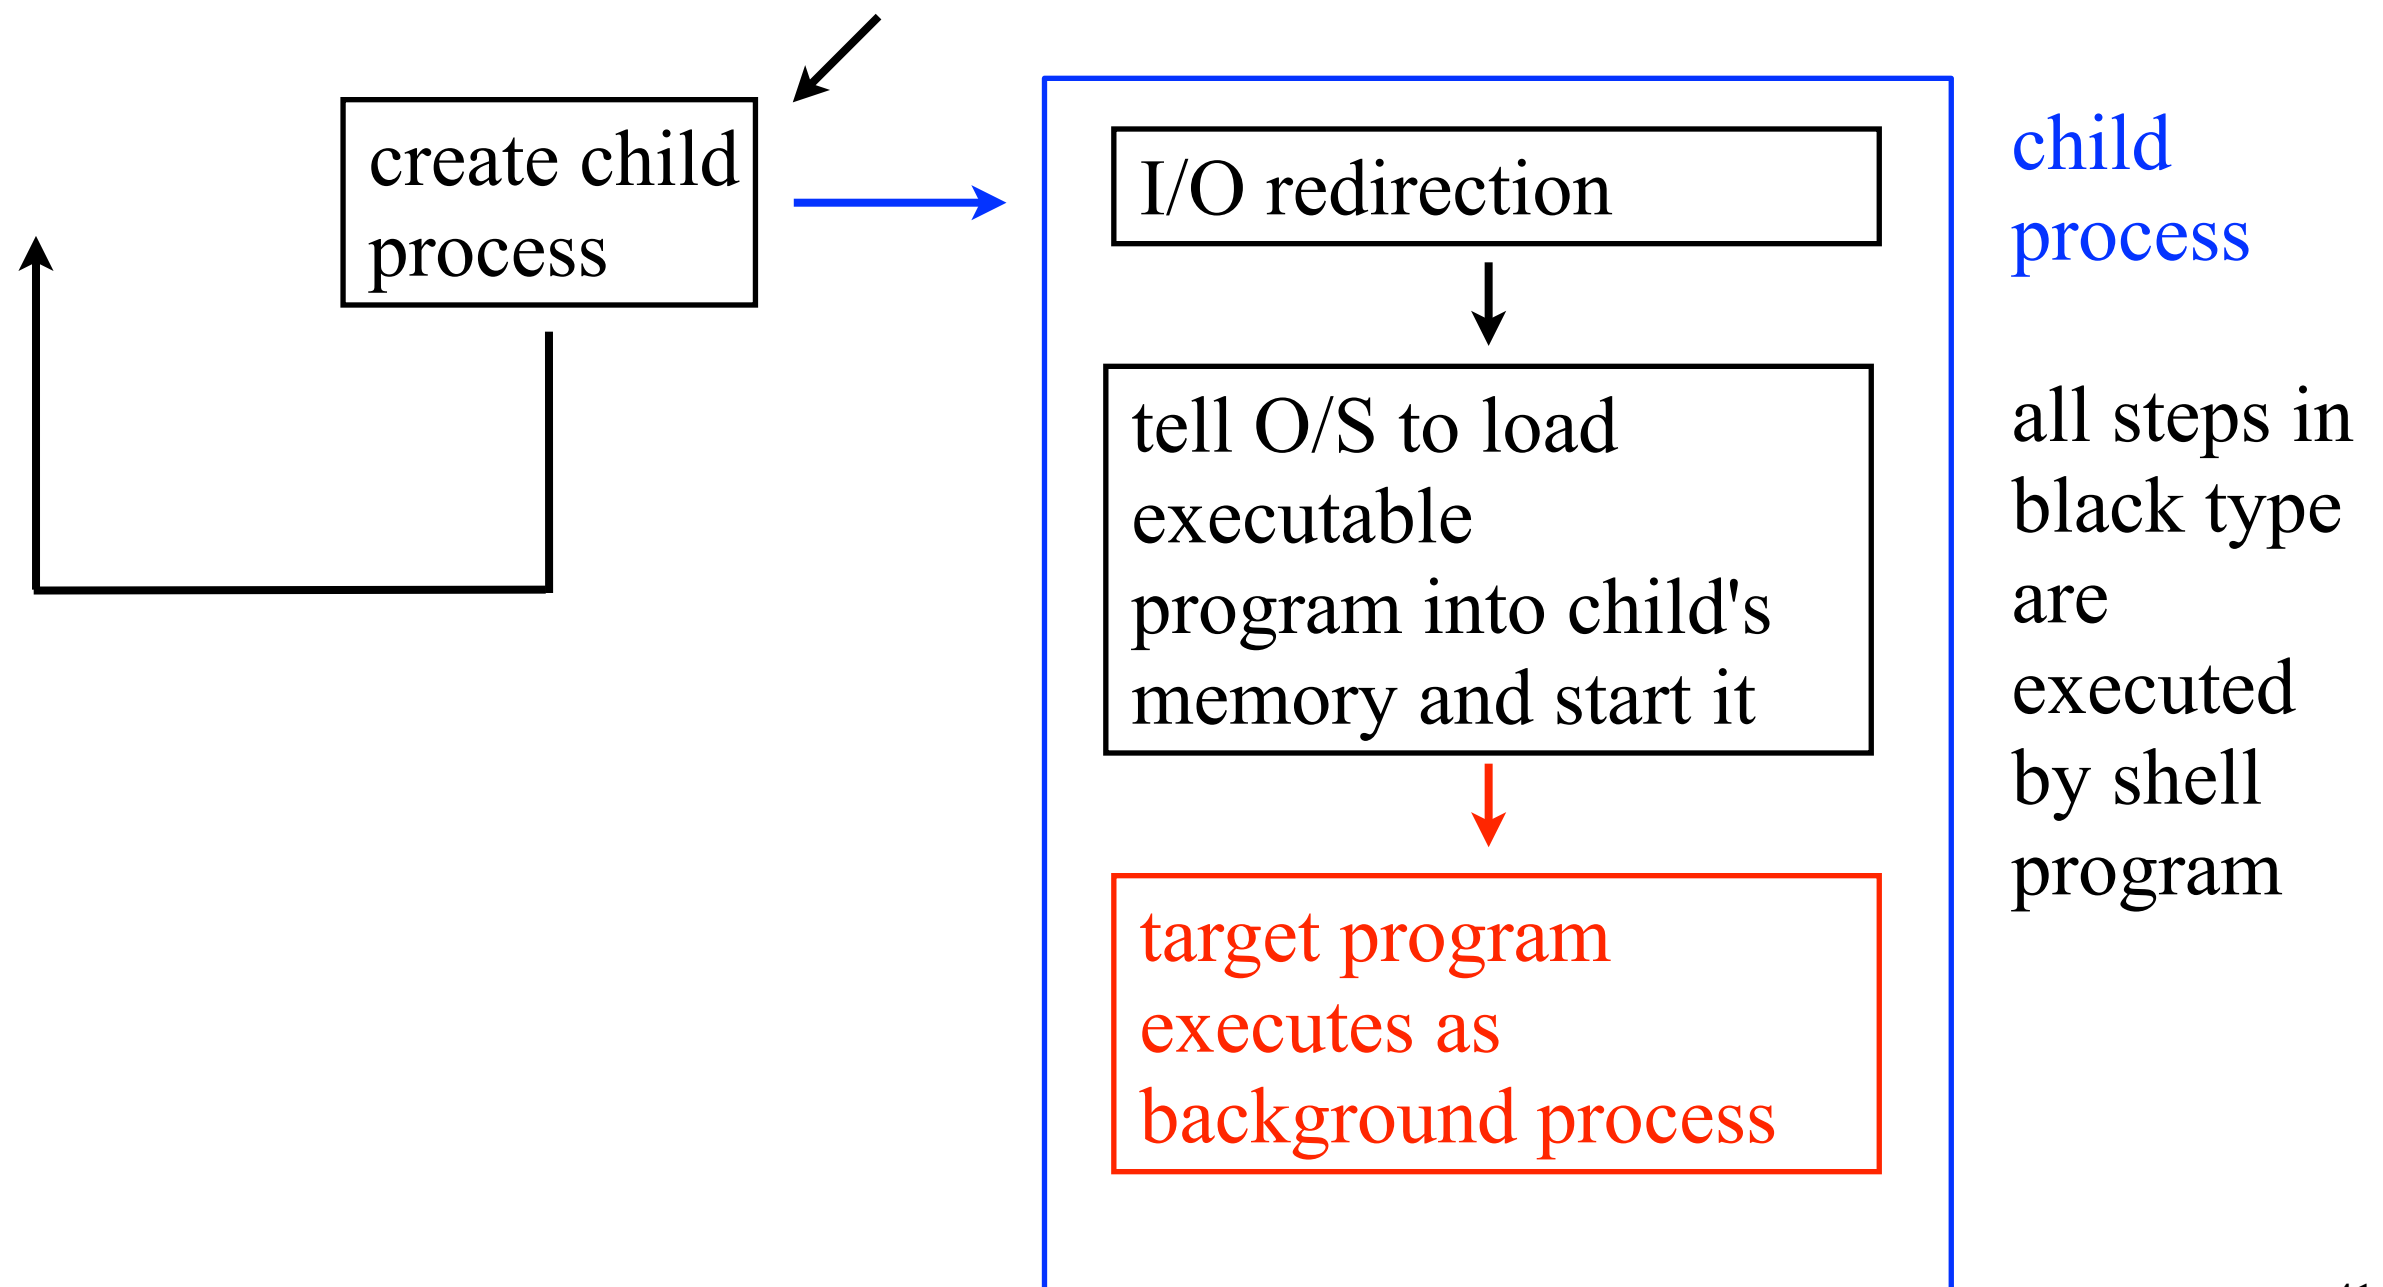
Basic Process Abstraction in UNIX

- ▶ process abstraction involved in executing a background command from the shell



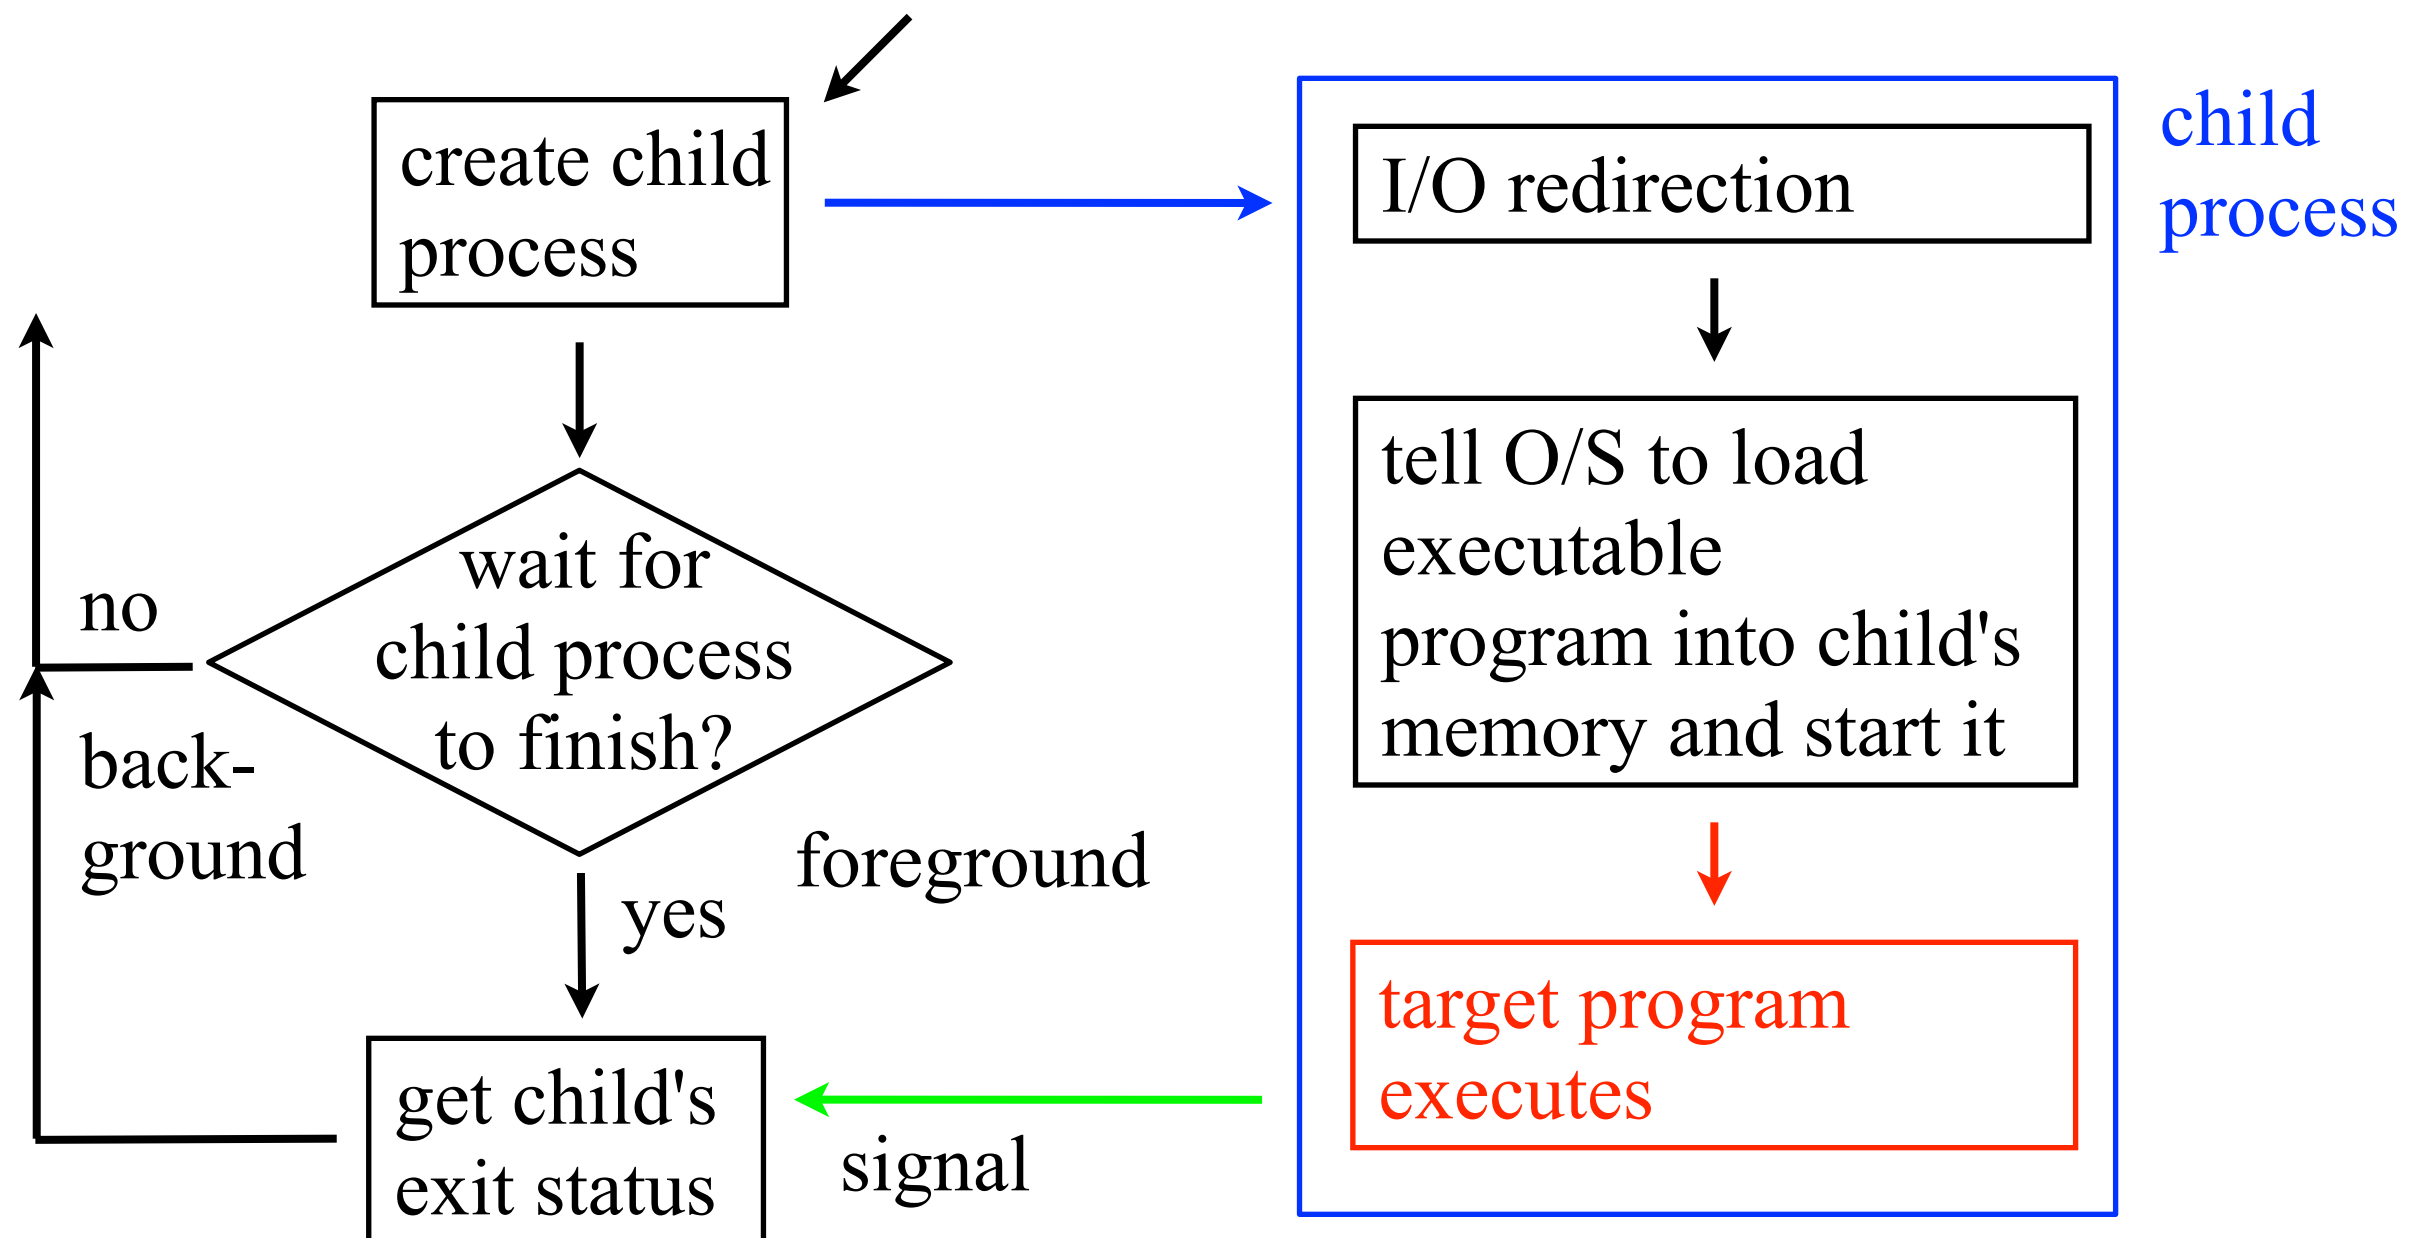
Basic Process Abstraction in UNIX

- ▶ process abstraction involved in executing a background command from the shell



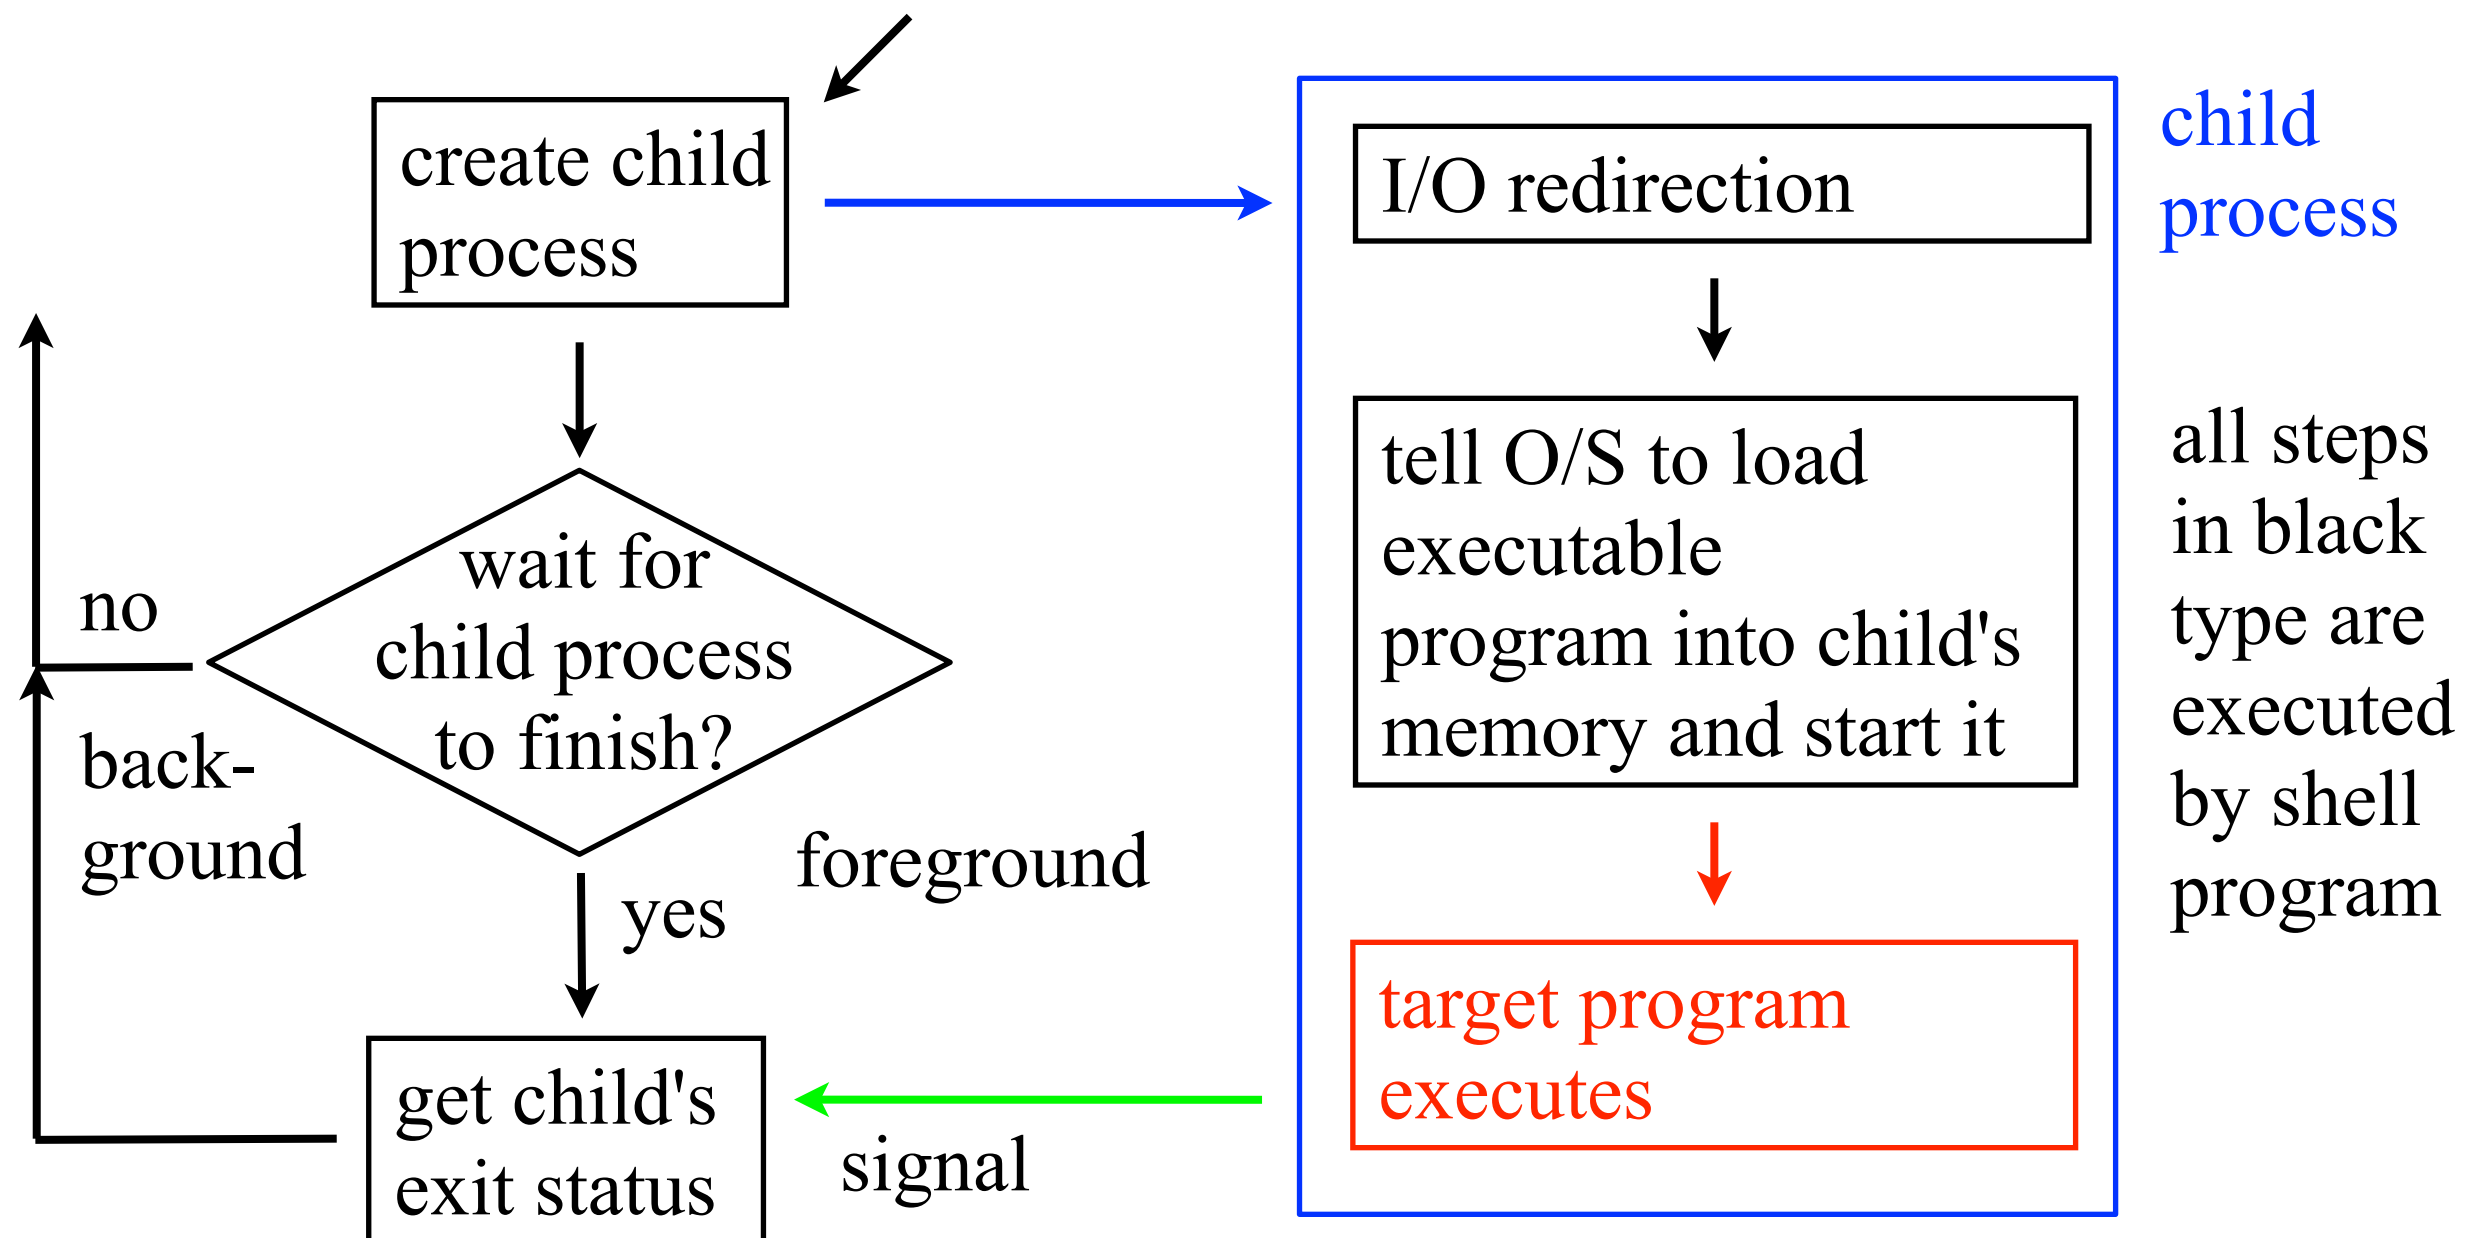
Basic Process Abstraction in UNIX

- ▶ process abstraction involved in executing a command from the shell



Basic Process Abstraction in UNIX

- ▶ process abstraction involved in executing a command from the shell



Processes and Jobs

- ▶ background process

- &
- bg

- ▶ suspend process:

- ^Z

- ▶ job:

- jobs
- %n

- ▶ foreground process:

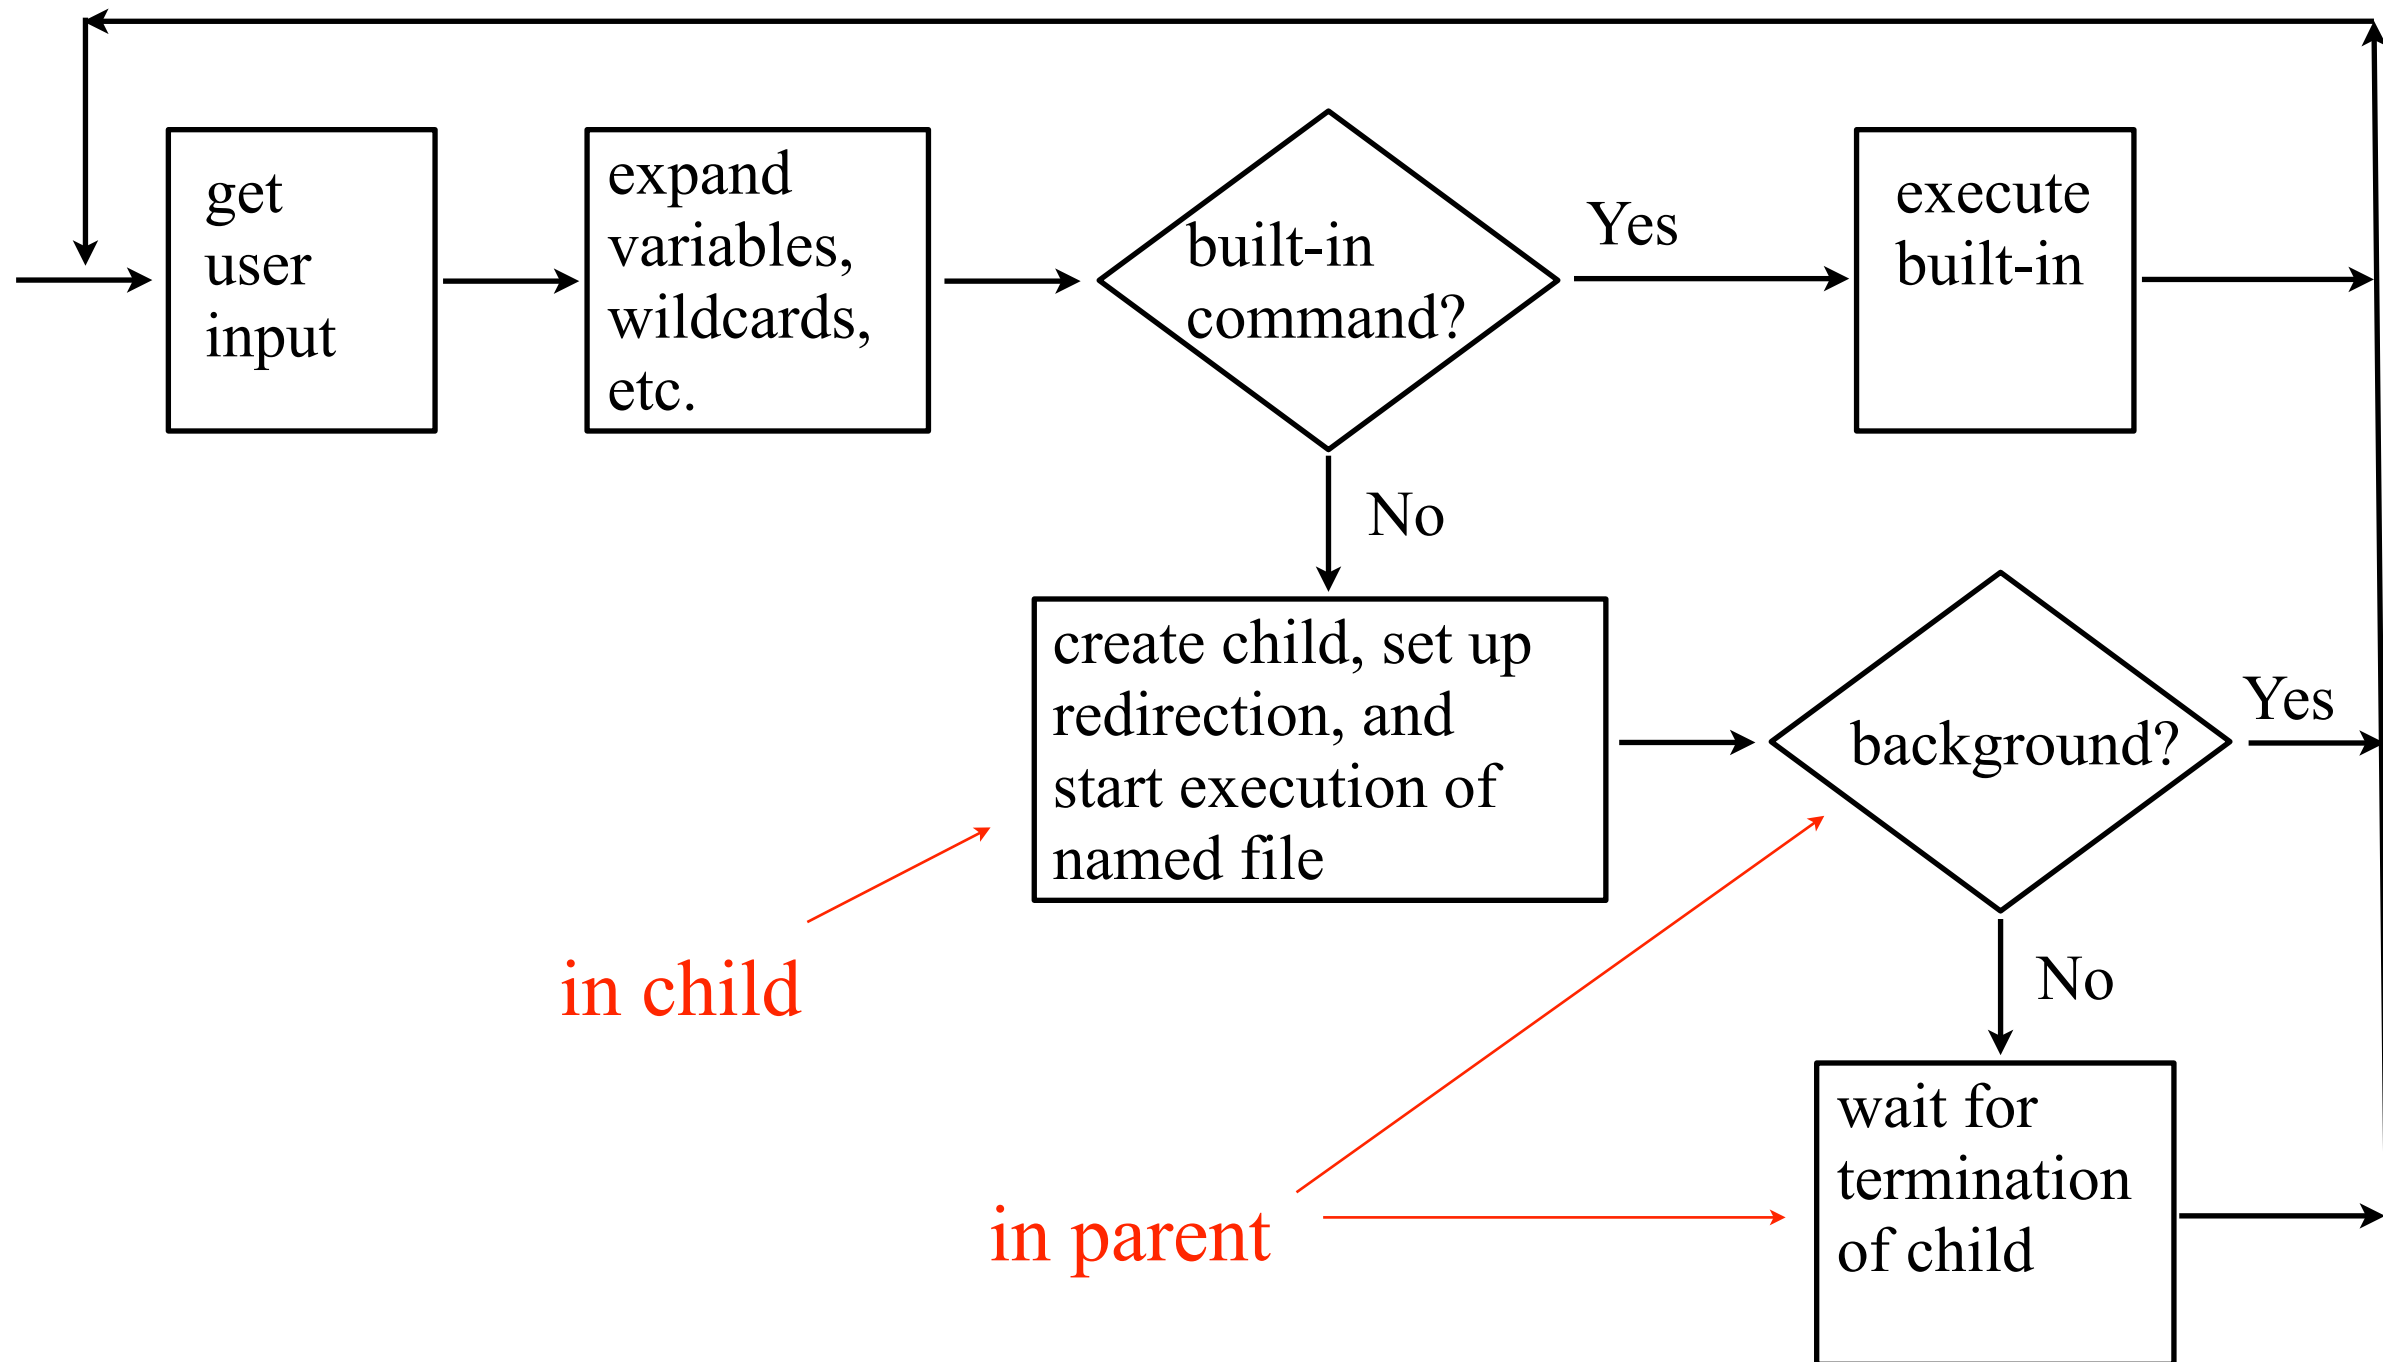
- fg %n

Processes and Jobs

► example involving `&`, `kill`, `uniq`

```
# /bin/bash
yes > raw_stuff &
kill -TERM %1
uniq -c raw_stuff
wc -l raw_stuff
rm raw_stuff
```

Basic Shell Operation



leave LINUX/UNIX shell ... for now